

# Workload-oriented Programming

Silviu S. Craciunas

Department of Computer Sciences  
University of Salzburg, Austria

-  
joint work with  
C. Kirsch, H. Röck, and A. Sokolova

July 2008

## Tiptoe

Prototype real-time operating system.

<http://tiptoe.cs.uni-salzburg.at>

- Processes are sequences of actions
- Actions may have workload parameters that determine the amount of work involved
- We are interested in the temporal performance of actions with respect to the workload involved

# Example

Consider a process that:

- allocates memory for a video stream
- reads the video stream from a network connection
- compresses the stream
- stores it on disk
- deallocates the used memory

# Example

## Pseudo-code

```
loop {  
  int number_of_frames=determine_rate();  
  allocate_memory(number_of_frames);  
  read_from_network(number_of_frames);  
  compress_data(number_of_frames);  
  write_to_disk(number_of_frames);  
  deallocate_memory(number_of_frames);  
} until (done);  $\leftarrow$  requirements
```

# Example

## Pseudo-code


```
loop {  
  int number_of_frames=determine_rate();  
  allocate_memory(number_of_frames);← requirements 1  
  read_from_network(number_of_frames);← requirements 2  
  compress_data(number_of_frames);← requirements 3  
  write_to_disk(number_of_frames);← requirements 4  
  deallocate_memory(number_of_frames);← requirements 5  
} until (done);
```

# Example


## Pseudo-code

```
loop {  
    int number_of_frames=determine_rate();  
    > allocate_memory(number_of_frames);  
    read_from_network(number_of_frames);  
    compress_data(number_of_frames);  
    write_to_disk(number_of_frames);  
    deallocate_memory(number_of_frames);  
} until (done);
```

# Response-time function




# Improving latency over throughput






# Improving throughput over latency



# Response-time function



# Response-time function

## Informally


The response-time (RT) function  $f_R$  characterizes the action's response time bound for a given workload, independently of any previous or concurrent actions.

## Formally


A response-time (RT) function is a discrete function

$$f_R : \mathbb{N} \rightarrow \mathbb{Q}^+$$

# Execution-time function



# Execution-time function



# Execution-time function

## Informally

The execution-time (ET) function  $f_E$  characterizes the action's execution time bound for a given workload, in the absence of any concurrent actions.

## Formally

An execution-time (ET) function is a discrete function

$$f_E : E_D \rightarrow \mathbb{Q}^+.$$

where  $E_D \subseteq \mathbb{N}$  is called the execution domain.

# Scheduling and Admission

## Process scheduling

How do we efficiently schedule processes on the level of individual process actions?

## Process admission

How do we efficiently test schedulability of newly arriving processes?


# Virtual periodic resources

## Virtual periodic resources

Period  $\pi$

Limit  $\lambda$

Utilization  $\frac{\lambda}{\pi}$



We use our modified version of the CBS<sup>1</sup> which allows different  $\pi$  and  $\lambda$  for a process.

<sup>1</sup>Giorgio Buttazzo and Enrico Bini - *Optimal Dimensioning of a Constant Bandwidth Server* in Proc. RTSS 2006




# Virtual periodic resources

- Each process declares a finite set of virtual periodic resources
- Each process action of a process uses exactly one virtual periodic resource declared by the process

# Scheduled response time

The scheduled response time  $f_S(w)$  is the time from the arrival of an action until it completes.



# Schedulability result

A set of processes is schedulable under EDF if


- For every action of every process,  $\lambda$  and  $\pi$  “sample”  $f_R$  such that

$$\forall w \in E_D, f_S(w) \leq f_R(w)$$

- The schedulability test holds for the set of processes

$$\sum_P \max_R \frac{\lambda_{PR}}{\pi_{PR}} \leq 1$$

# Response time sampling




# Tiptoe compositionality

$$\forall f_S, f_S', \forall w \in E_D$$

$$0 \leq |f_S(w) - f_S'(w)| \leq \pi_a - 1$$

if the schedulability test holds




## Response-time jitter

Any individual action  $a$  of a process maintains its response time within  $\pi_a - 1$  even when adding/removing concurrent processes.


# Implementation and performance

	list	array	matrix
time	$O(n^2)$	$O(\log(t) + n \cdot \log(t))$	$\Theta(t)$
space	$\Theta(n)$	$\Theta(t + n)$	$\Theta(t^2 + n)$



# Implementation and performance


	list	array	matrix
time	$O(n^2)$	$O(\log(t) + n \cdot \log(t))$	$\Theta(t)$
space	$\Theta(n)$	$\Theta(t + n)$	$\Theta(t^2 + n)$






# Implementation and performance

	list	array	matrix
time	$O(n^2)$	$O(\log(t) + n \cdot \log(t))$	$\Theta(t)$
space	$\Theta(n)$	$\Theta(t + n)$	$\Theta(t^2 + n)$



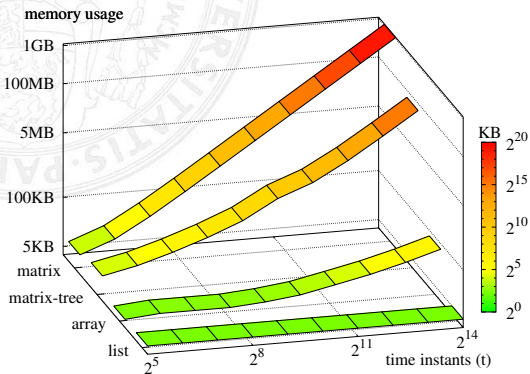
# Implementation and performance

	list	array	matrix
time	$O(n^2)$	$O(\log(t) + n \cdot \log(t))$	$\Theta(t)$
space	$\Theta(n)$	$\Theta(t + n)$	$\Theta(t^2 + n)$



# Implementation and performance

	list	array	matrix
time	$O(n^2)$	$O(\log(t) + n \cdot \log(t))$	$\Theta(t)$
space	$\Theta(n)$	$\Theta(t + n)$	$\Theta(t^2 + n)$



# Conclusion

The model enables

- trading off throughput and latency
- sequential and concurrent process composition
- controlling the response-time variance (jitter) of individual process actions

# Conclusion

The model enables


- trading off throughput and latency
- sequential and concurrent process composition
- controlling the response-time variance (jitter) of individual process actions

# Conclusion

The model enables

- trading off throughput and latency
- sequential and concurrent process composition
- controlling the response-time variance (jitter) of individual process actions

## Fly the JAviator with Tiptoe



## Homepage

<http://javiator.cs.uni-salzburg.at>



Thank you!



# Scheduled response time

$\forall w \in U_D, f_S(w) \leq f_R(w) + \pi$  if

- $\frac{\lambda}{\pi} = c_U$
- $\pi$  divides  $f_R(w)$  evenly
- $\sum_P \max_R \frac{\lambda_{PR}}{\pi_{PR}} \leq 1$

# Scheduled response time

$\forall w \in U_D, f_S(w) \leq f_R(w)$  if

- $\frac{\lambda}{\pi} = c_U$
- $0 < \pi \leq d_R - \frac{d_E}{c_U}$
- $\pi$  divides  $d_R$  and  $f_R(w) - d_R$  evenly
- $\sum_P \max_R \frac{\lambda_{PR}}{\pi_{PR}} \leq 1$