



RTNS 2024

TTTech

Taming System Complexity

Global Scheduling and the Time-Triggered Approach

Silviu S. Craciunas

Principal scientist @ TTTech Labs



TTTech project highlights

TTTech



TTTech Auto safety and robustness in series production

Hardware development until the A-Sample grade



Scalable distributed control system from TTTech Industrial into Vestas 2, 4 and 6 Megawatt turbines



Collins Aerospace Power System and Air Management for the Boeing 787-8, 787-9, 787-10 family



TTTech TT Ethernet Switch and End System IP products for the avionics system.

TT Ethernet is the “nervous system” i.e. avionics network platform of the space craft

TTTech Group Key facts



Founded in 1998, headquartered in Vienna, Austria, with 19 offices in 14 countries worldwide



Products in 1173 production programs



Connected companies: TTTech Auto, TTTech Industrial, TTControl, RT-RK

2,300+

Employees/
subcontractors



36

ongoing
research
projects



10,000

Vestas turbines



1.5 billion

operating hours of
railway systems



>2 million

kilometers covered by
TTEthernet in deep space

60

Nations represented
by our workforce

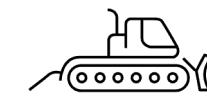


635,000 ECUs &
82,000 displays



3 million

cars with MotionWise
on the road



1.9 million

kilometers of slopes



1 billion

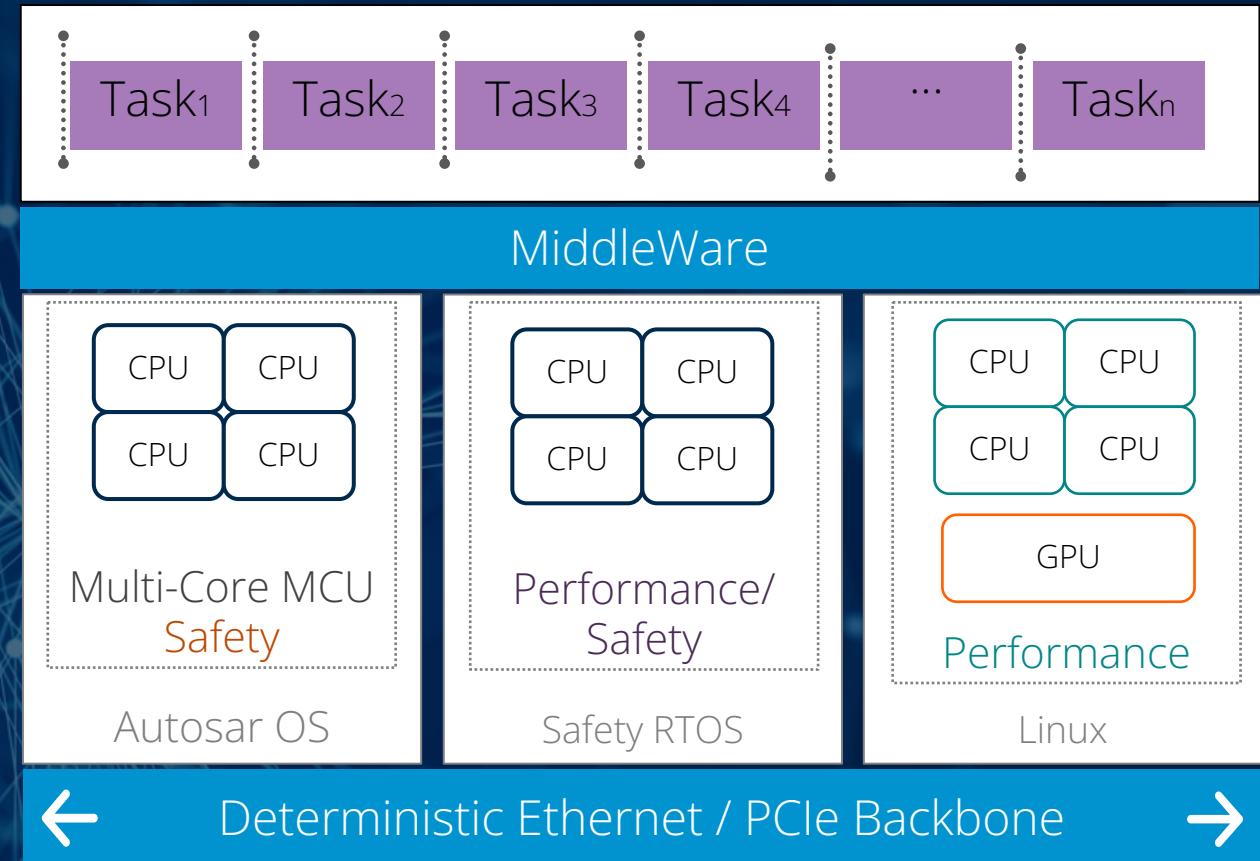
flight hours

Numbers as of March 2022

Systems and platform



Heterogeneous multi-core multi-SoC platforms featuring a variety of CPUs and GPUs running at different speeds, which are interconnected through either a deterministic Ethernet backbone (TSN) or through PCIe



Large number of safety-critical tasks with increasing code size that have complex timing requirements beyond just deadlines

Cause-effect chains

Cause-effect chains are a characteristic of automotive software:

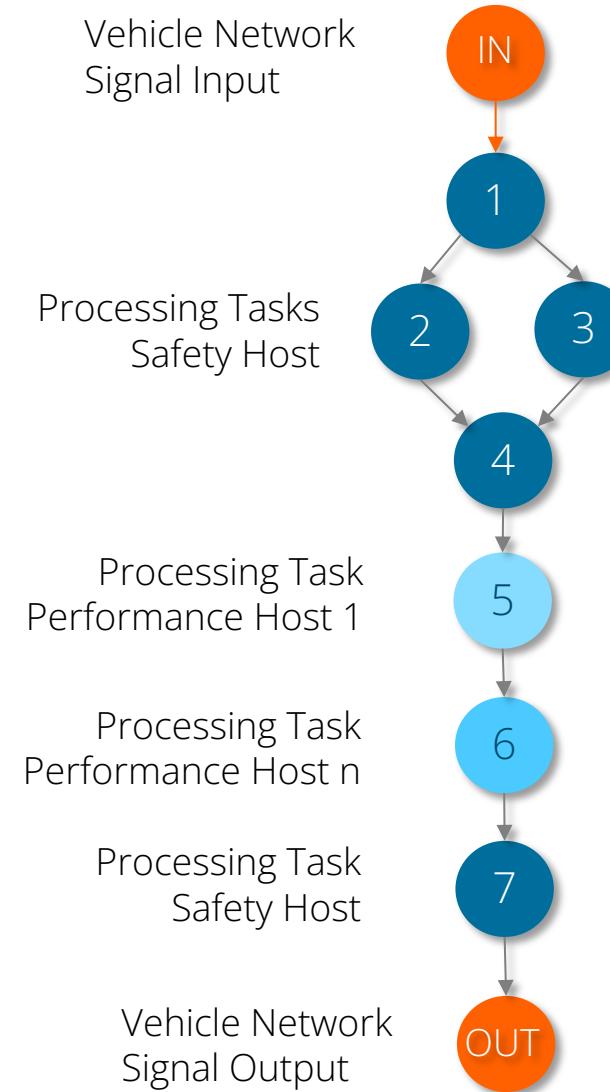
- provide additional timing and dependency requirements on the execution of tasks
- can span across multiple activation patterns
- include multiple tasks, even the same task multiple times
- have priorities and end-to-end latencies
- include communication latencies
- task execution and network communication must be aligned across SoCs and CPUs

Semantics:

- Reaction delay
- Data age
- Order

ADAS Domain

Vehicle Network
Signal Input



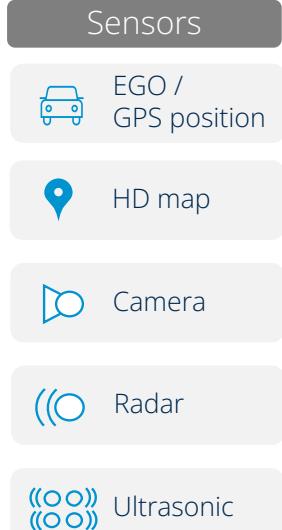
Processing Tasks
Safety Host

Processing Task
Performance Host 1

Processing Task
Performance Host n

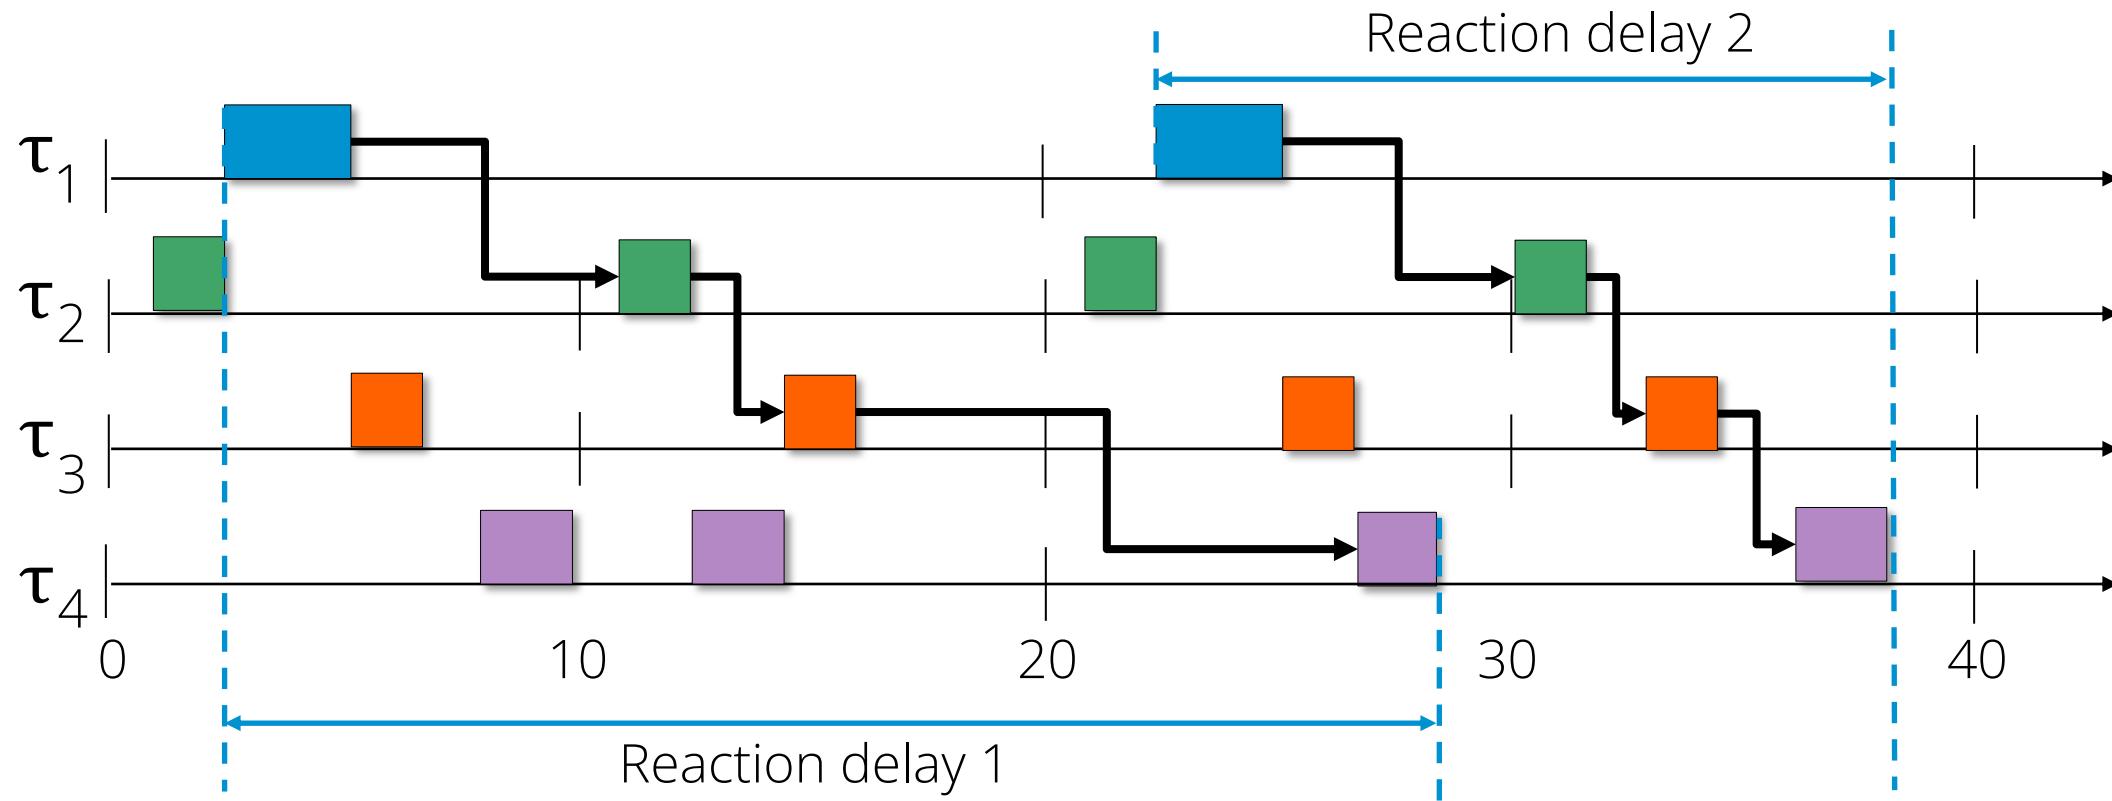
Processing Task
Safety Host

Vehicle Network
Signal Output



Cause-effect chains: reaction delay

After each event, there must be a reaction within the maximum allowed time

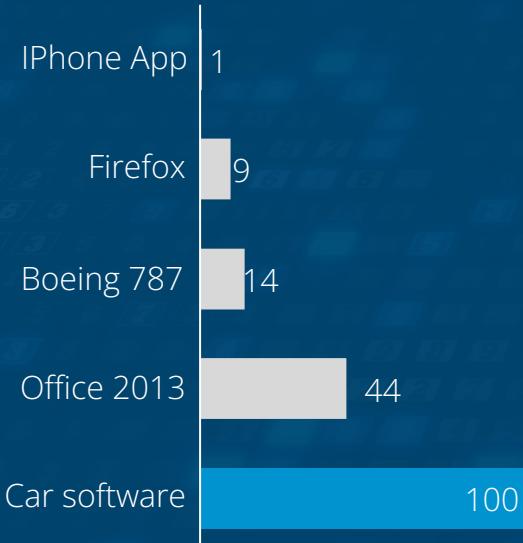


An aerial photograph of a multi-level highway interchange at night. The image is dominated by the warm orange glow of streetlights and vehicle headlights, which create long, streaky light trails against the dark sky. The highway structures are illuminated with blue and white lights, forming a complex network of intersecting beams. The perspective is from above, looking down the length of the roads.

Complexity

System complexity – ADAS example

Increasing complexity of car software (lines of code, millions)



Increasing complexity of timing requirements

- Critical event/data-driven chains
- Hard real-time and determinism requirements
- Orchestration needs of heterogeneously distributed applications
- Mixed-criticality and Freedom From Interference requirements
- Complex task dependencies



Design-time guarantees

- Testing-based approaches are not sufficient anymore
- Correct by design approach with mathematically-proven guarantees
- This implies more knowledge (e.g., timing budgets) and solutions to complex scheduling/configuration problems

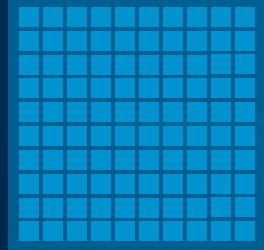
"Design approaches not based on mathematically proven real time scheduling properties are prone to missing deadlines during unusual operational conditions" UL4600 Standard for Safety for the Evaluation of Autonomous Products



Giving design-time guarantees implies solving the configuration/scheduling problem

Complexity in numbers

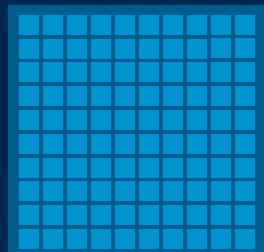
Software demands



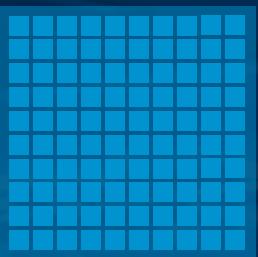
100 Applications



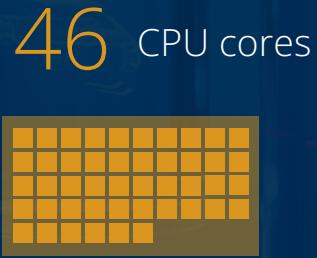
400 Tasks



400 Tasks



Computation unit resources



46 CPU cores



6 Hardware accelerators

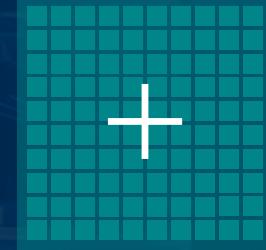


2 TSN switches

1024 virtual links



System constraints



100+ time, latency, precedence constraints etc.

Solution space

10⁵⁰⁰⁰

10⁵

possible configurations

correct configurations



Atoms in the universe
10⁸⁰

How can we solve anything given this complexity?

Classical (FP/EDF) approaches

Classical approaches can give design-time guarantees, sometimes even very efficiently.

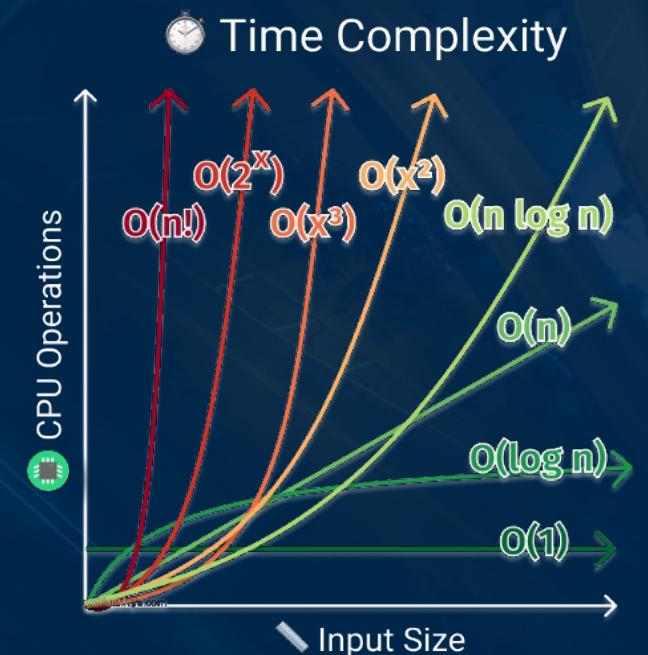
Deadline/rate monotonic, EDF, etc., are efficient ways to assign priorities, **but only for specific cases**:

- RM is efficient if: release jitter is 0, deadlines are equal to the periods, no dependencies
- DM is efficient if: release jitter is 0, deadlines are less than or equal to periods, no dependencies

Current automotive applications (e.g., ADAS/AD) are more complex:

- providing design-time guarantees implies checking correctness over all $n!$ priority orderings
- e.g., 25 tasks → 15511210043330985984000000 schedulability checks

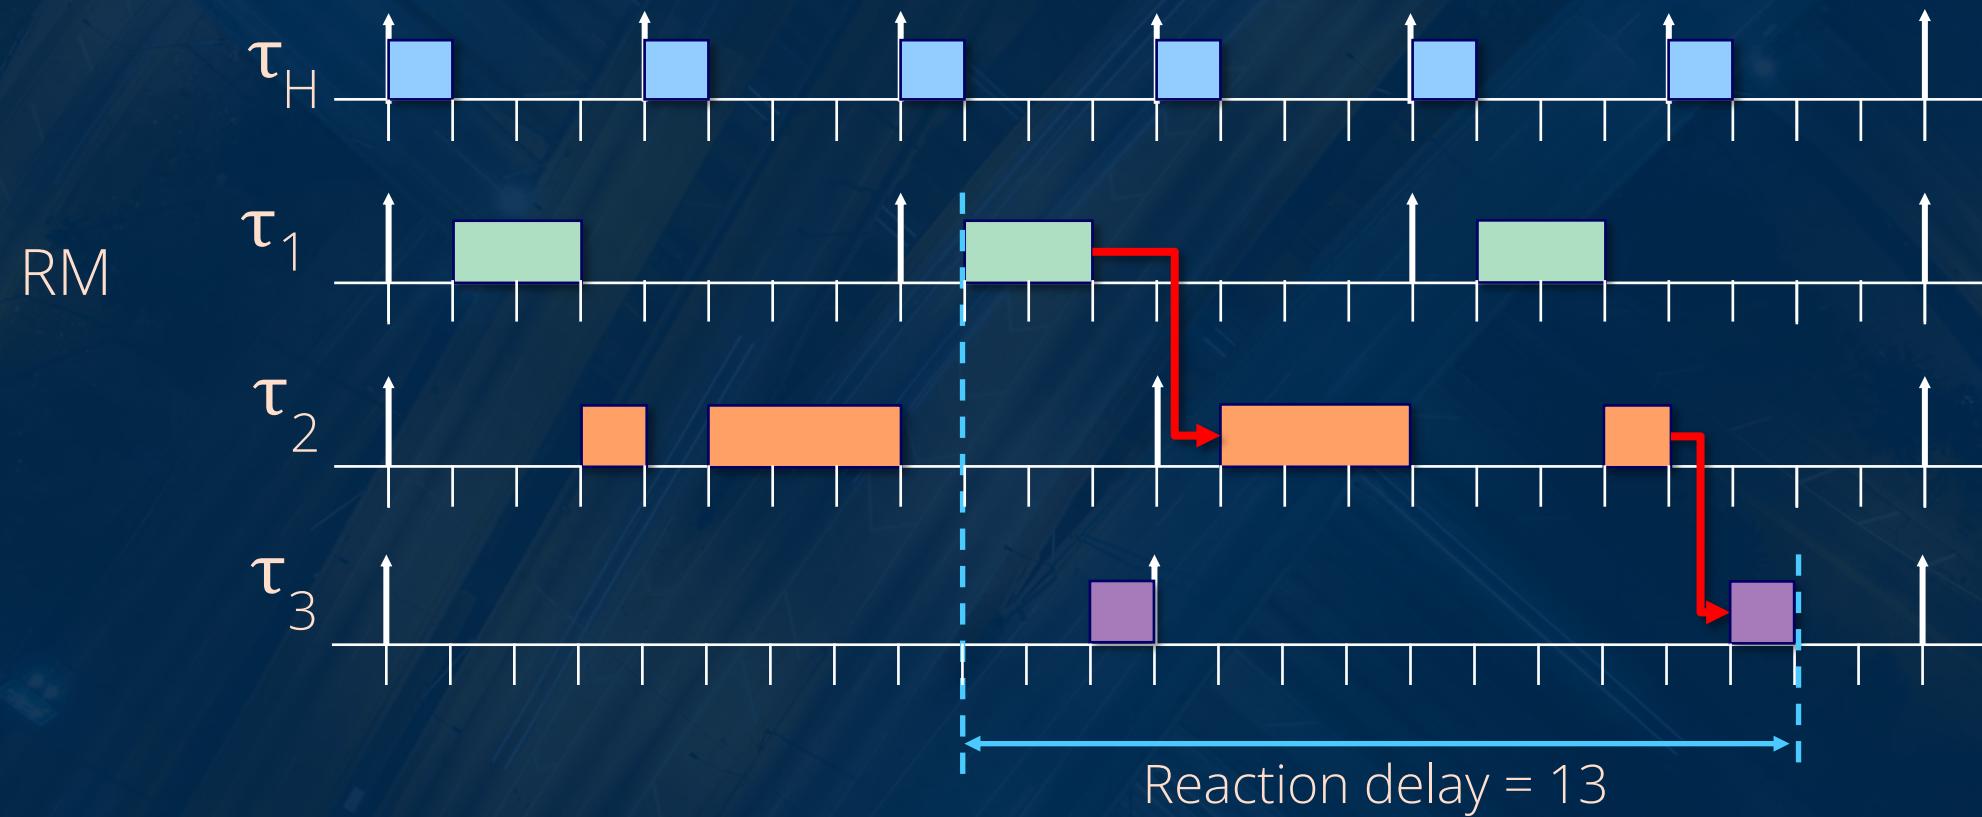
Giving design-time guarantees beyond deadlines is a very complex problem to solve for classical approaches.



Fixed-priority/EDF scheduling wrt. dependency chains

High-priority task: $T_H(1,4)$

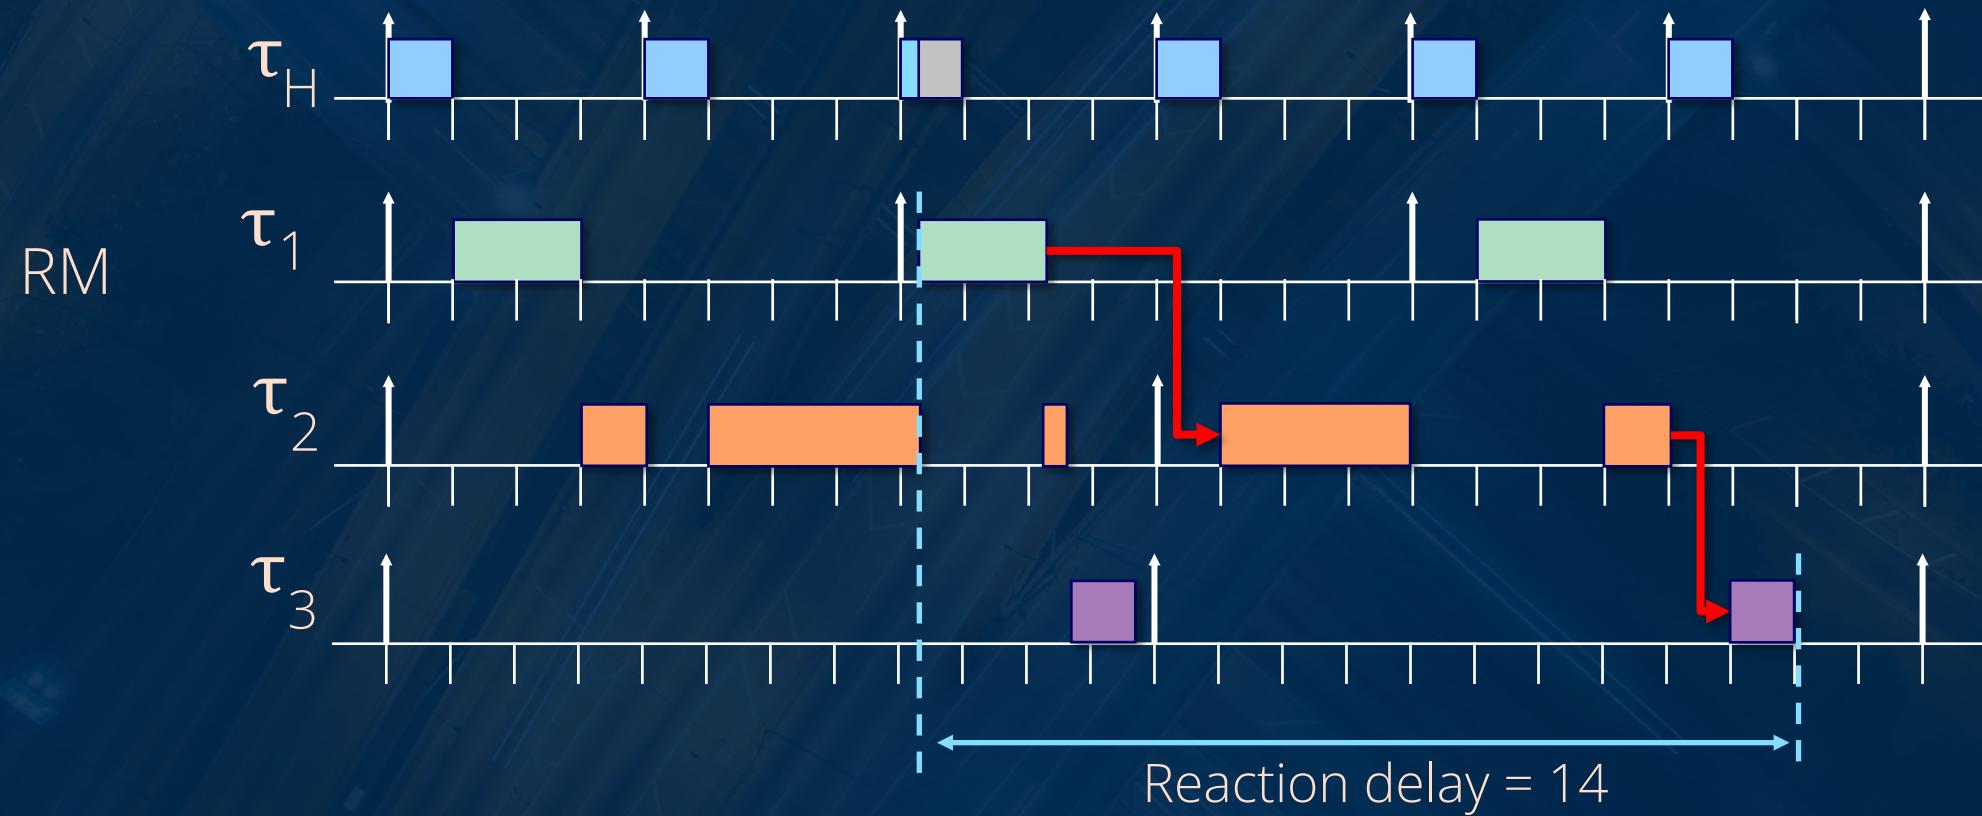
Chain: $T_1(2,8), T_2(4,12), T_3(2,12)$



Fixed-priority/EDF scheduling wrt. dependency chains

High-priority task: $T_H(1,4)$

Chain: $T_1(2,8), T_2(4,12), T_3(2,12)$



Complexity of runtime states

WCET is used for generating schedules and testing schedule correctness (also for FP)

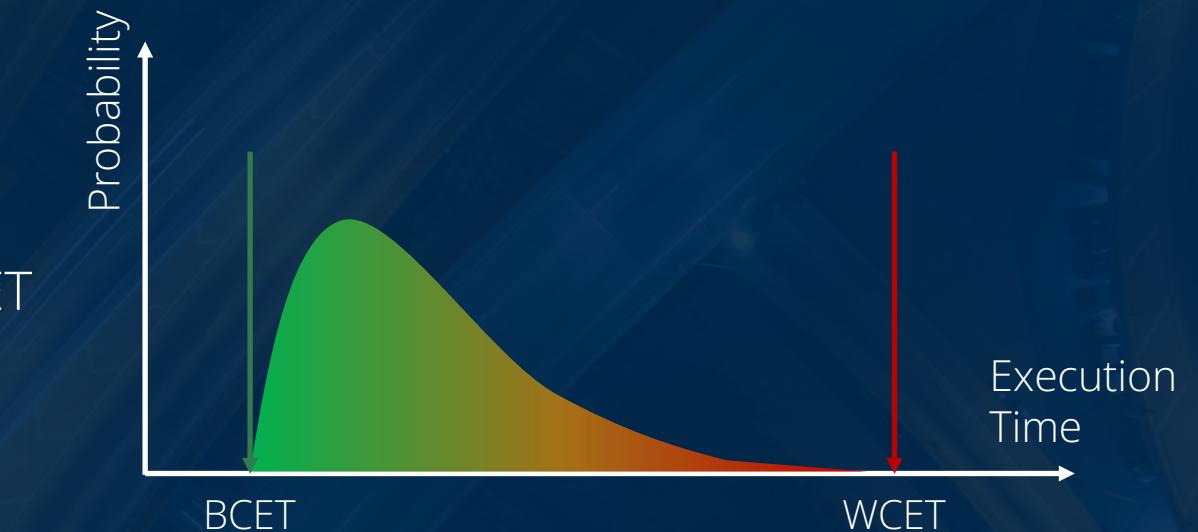
At runtime, tasks execute for less than the WCET

Example:

BCET = 1ms

WCET = 2ms

Microtick = 1us



1000 possible execution patterns

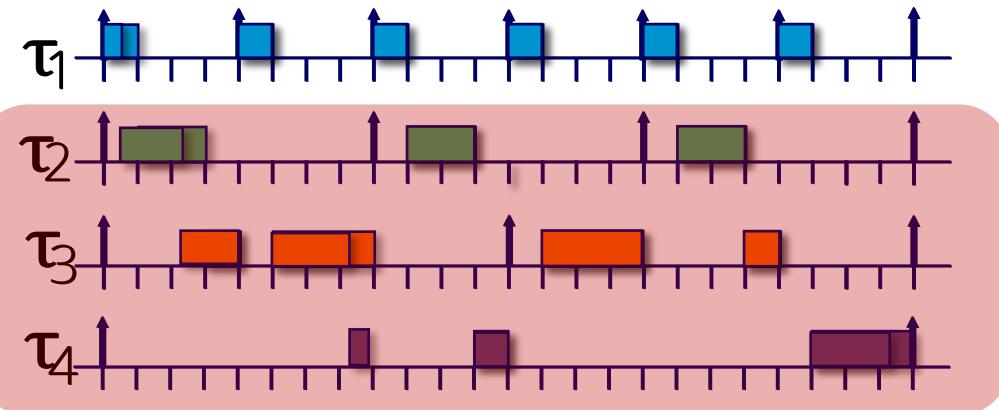
What happens if one task executes for less than the computed WCET?

Complexity of runtime states

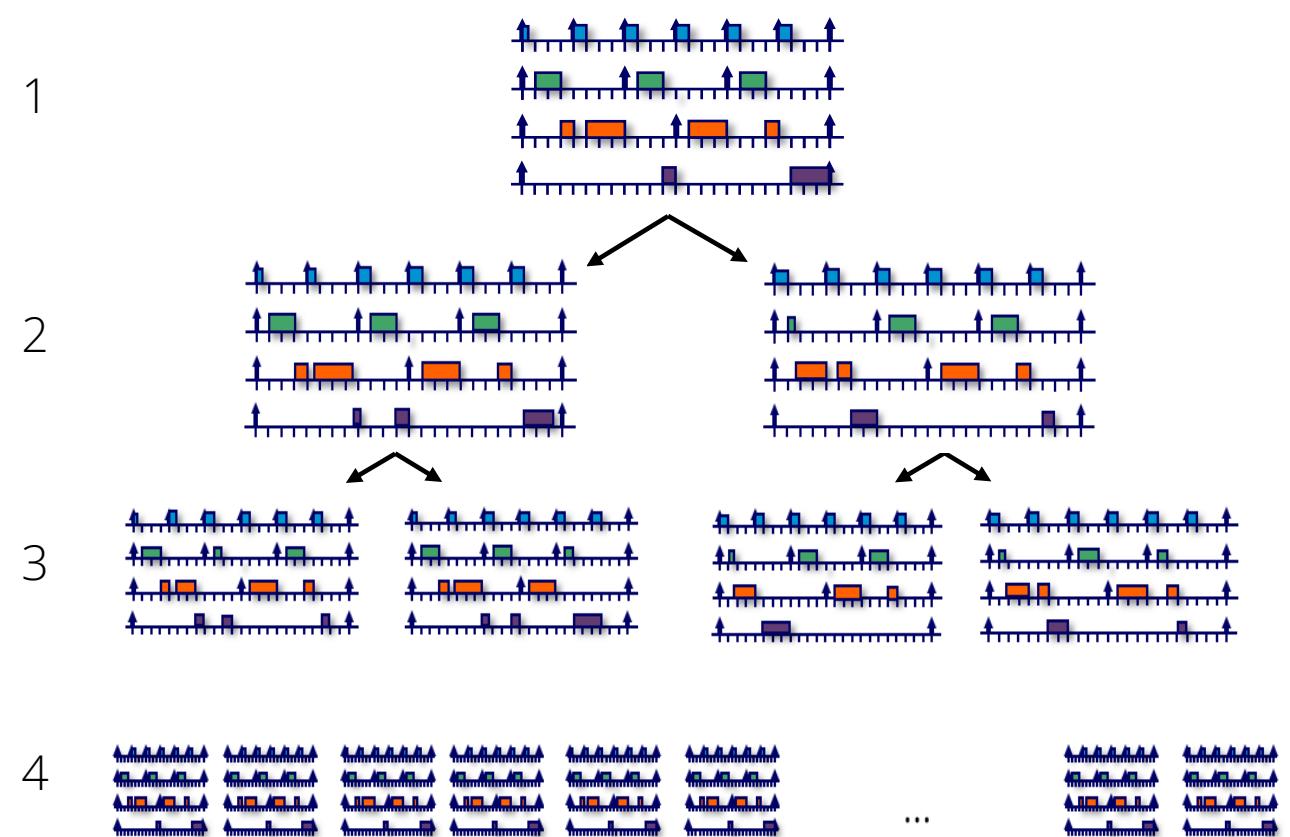
What happens if at some point in time, one task executes for less than the computed WCET?

Classical mechanisms

- changes the execution of all other tasks
- state change for the entire system



Non-TT (FP/EDF)

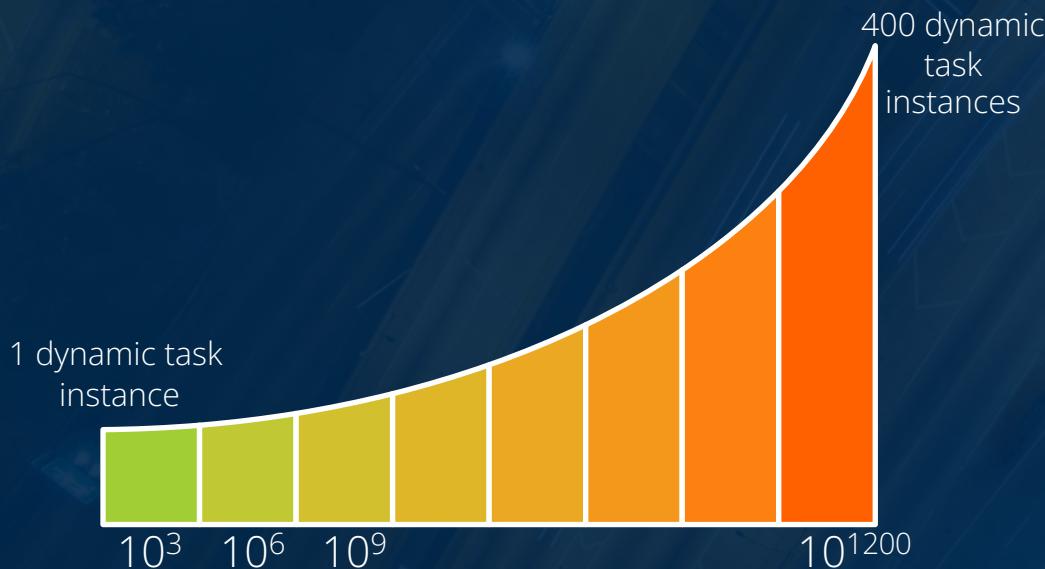


State space explosion

How many system states do I have to potentially check/look at to give design-time guarantees?

Non-time-triggered (FP/EDF/etc.):

- changes the execution of all other tasks
- state change for the entire system



Number of possible assignments of tasks to cores



Number of priority orderings



Number of system states

$$52^{400} \times 400! \times 10^{1200} = 10^{2755}$$

The background of the image is a photograph of a bridge at night. The bridge curves from the bottom left towards the center, with its surface appearing as a series of bright, glowing orange and yellow streaks due to long exposure photography. Below the bridge is a dark blue lake. In the distance, there are dark, silhouetted mountains under a sky that transitions from a pale yellow near the horizon to a deep blue at the top.

Time-Triggered Scheduling

The Time-Triggered advantage

Time-Triggered provides correct-by-design real-time guarantees!

Advantages of the Time-Triggered approach:

- **Complex timing requirements:** cause-effect chains, different types of jitter
- **Temporal isolation:** no starvation possible, temporal isolation between all tasks
- **Determinism/Predictability:** increased stability and testability, no unwanted run-time effects, many system properties become predictable, e.g., locks, task pre-emption
- **Re-simulation:** Equivalent behaviour between cloud and embedded target
- **Synchronization to communication:** stable real-time behaviour of cause-effect chains
- **Compositionality/Integration:** adding or modifying tasks can be done incrementally, system integration of SWCs from different suppliers is easier
- **Schedulability:** more correct configurations can be realized



For a more in-depth comparison between FPS and TT please see [Xu2000]

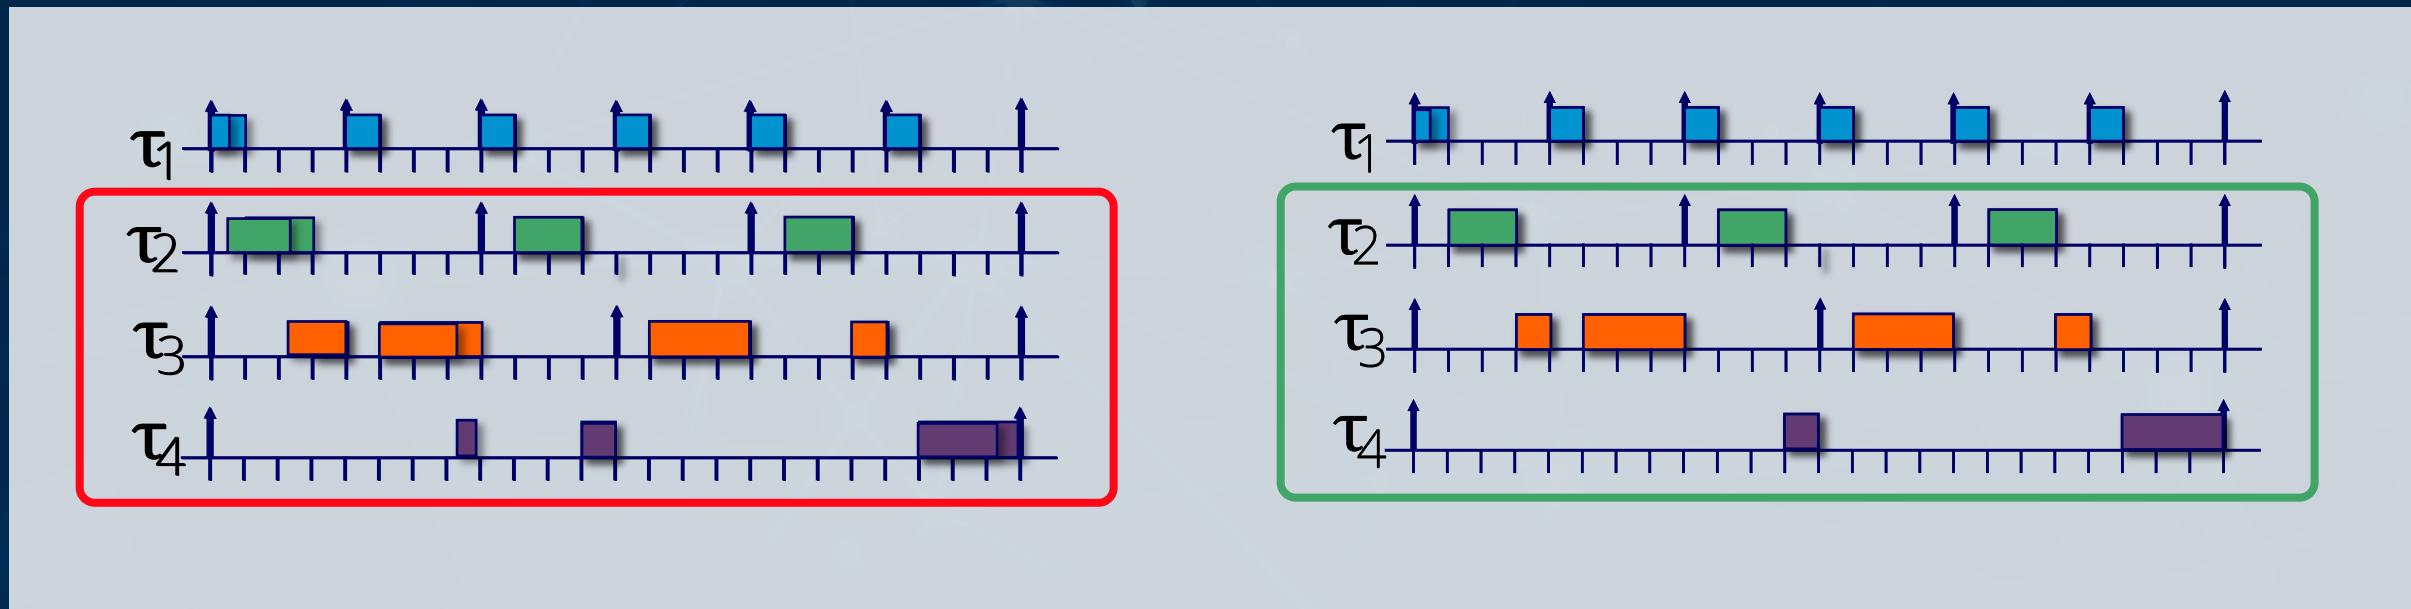
What happens if at some point in time, one task executes for less than the computed timing budget?

Non-TT (FP/EDF)

- changes the execution of all other tasks
- state change for the entire system

Time-Triggered

- execution of all other tasks remains the same
- no state change

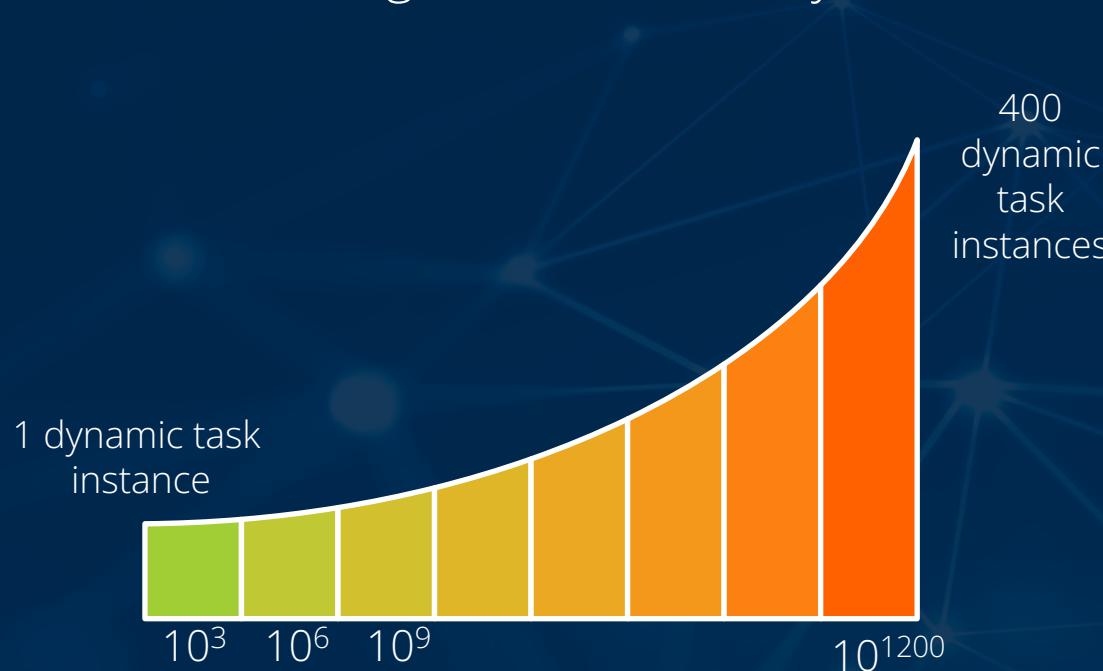


State space explosion

How many system states do I have to potentially check/validate?

Non-time-triggered (FP/EDF/etc.):

- changes the execution of all other tasks
- state change for the entire system



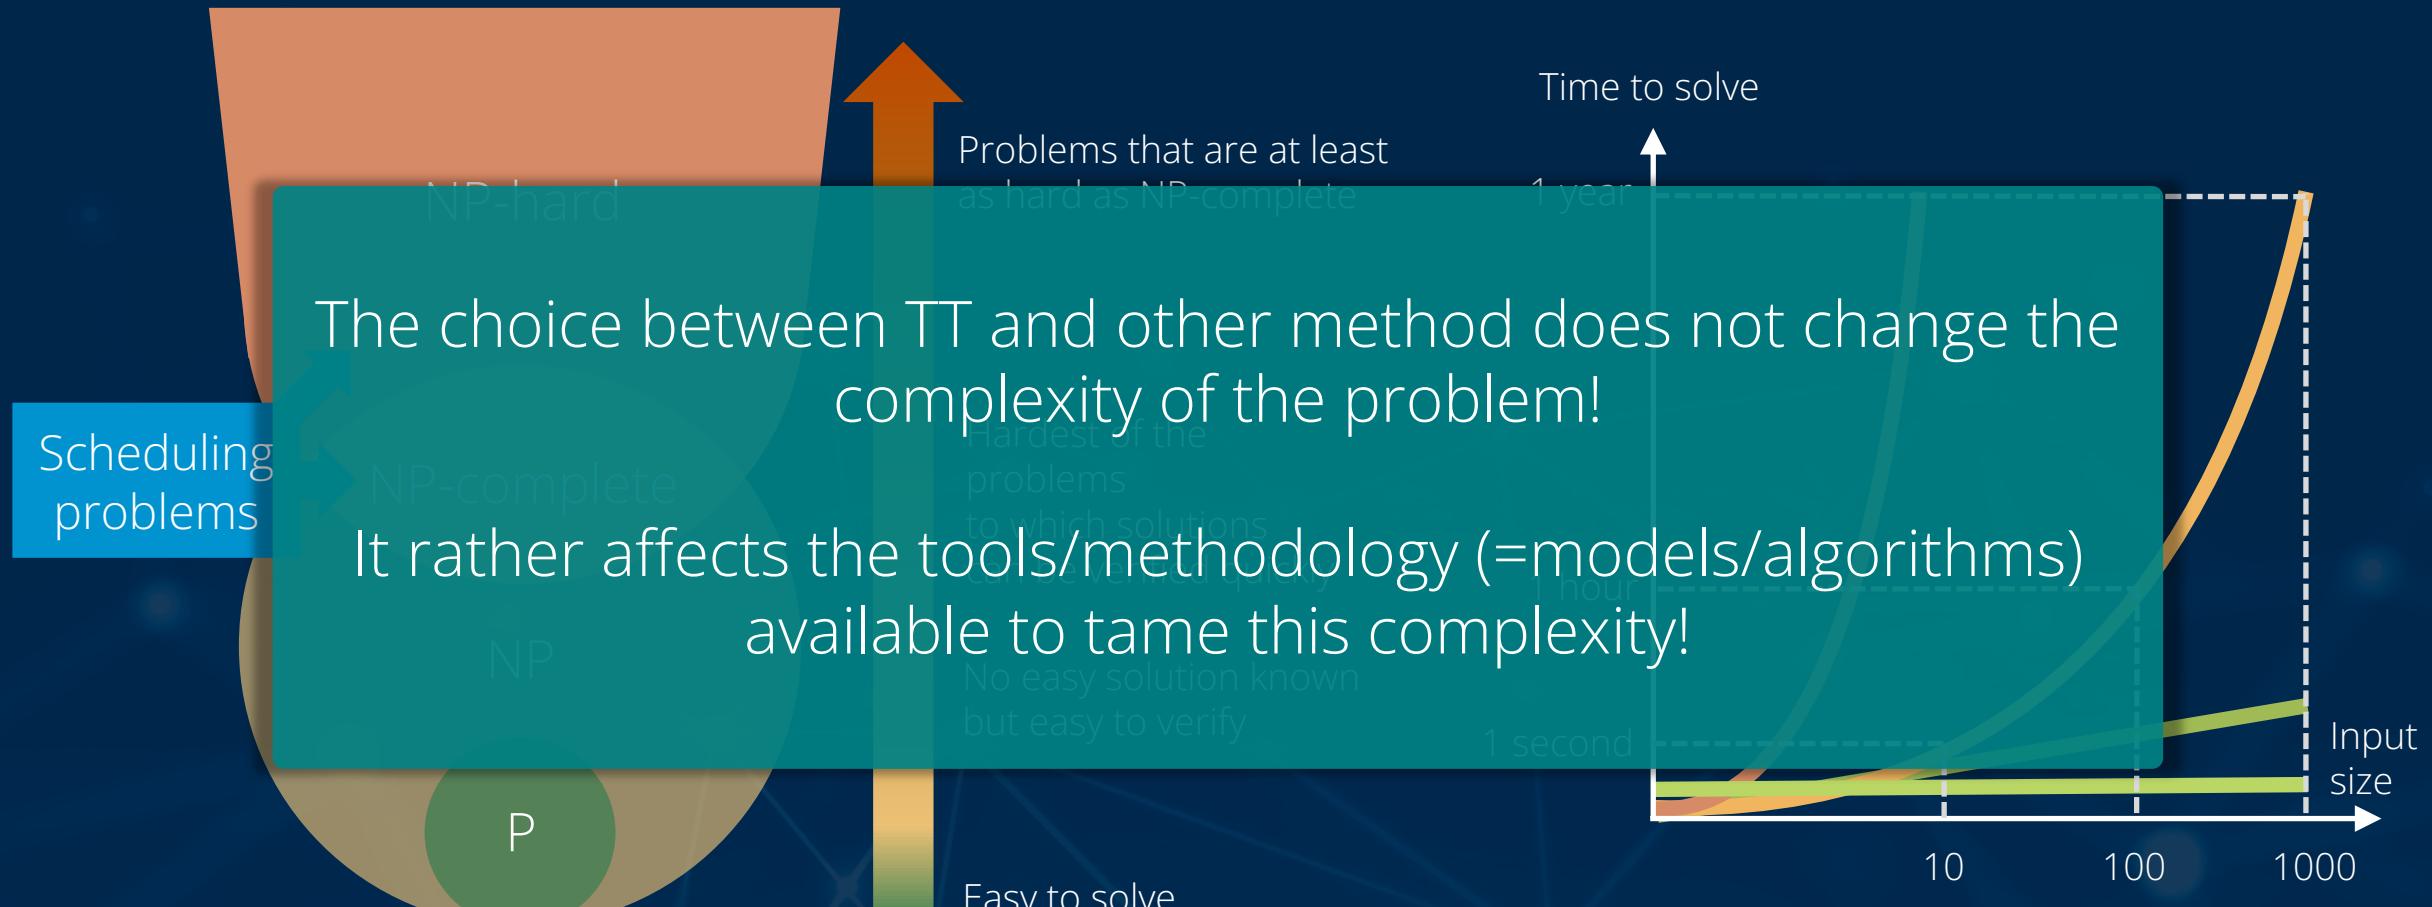
Time-triggered:

- execution of all other tasks remains the same
- no state change



BUT: Finding the static schedule table is still a very complex problem to solve

Complexity classes



Complexity is not always equal to algorithm efficiency: compare 1.00001^n vs n^{50}

Finding a needle in a haystack the size of the universe

TTTech





Taming complexity

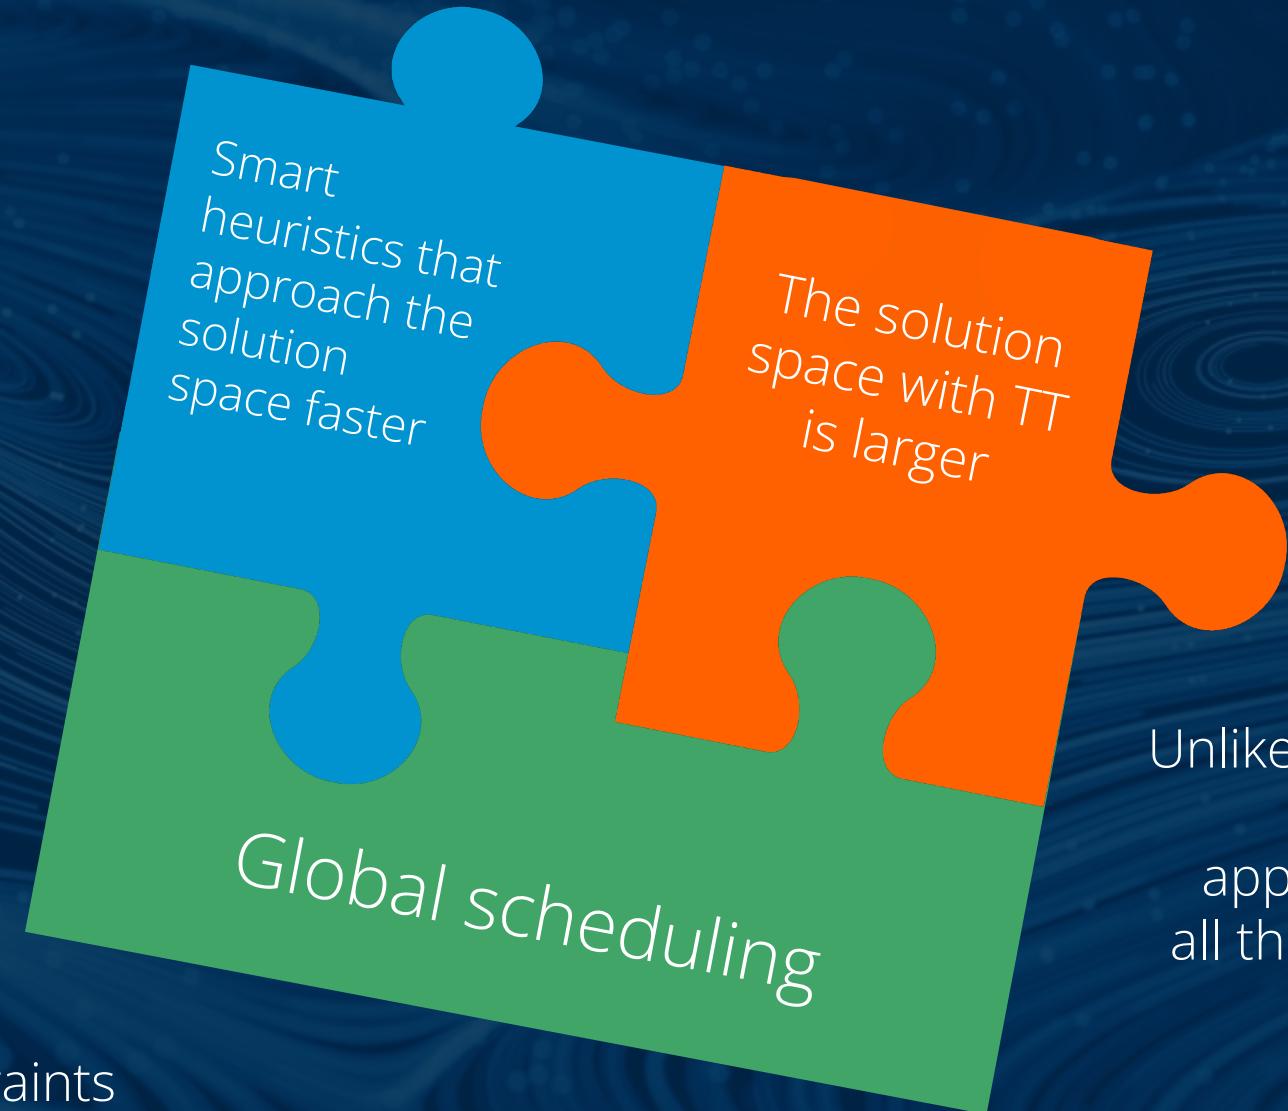
How do we tame the complexity



Heuristic algorithms that significantly increase the probability of finding correct schedules



Global scheduling that has a holistic view of the system and constraints



Unlike other techniques, the time-triggered approach can express all the correct solutions (solution space)

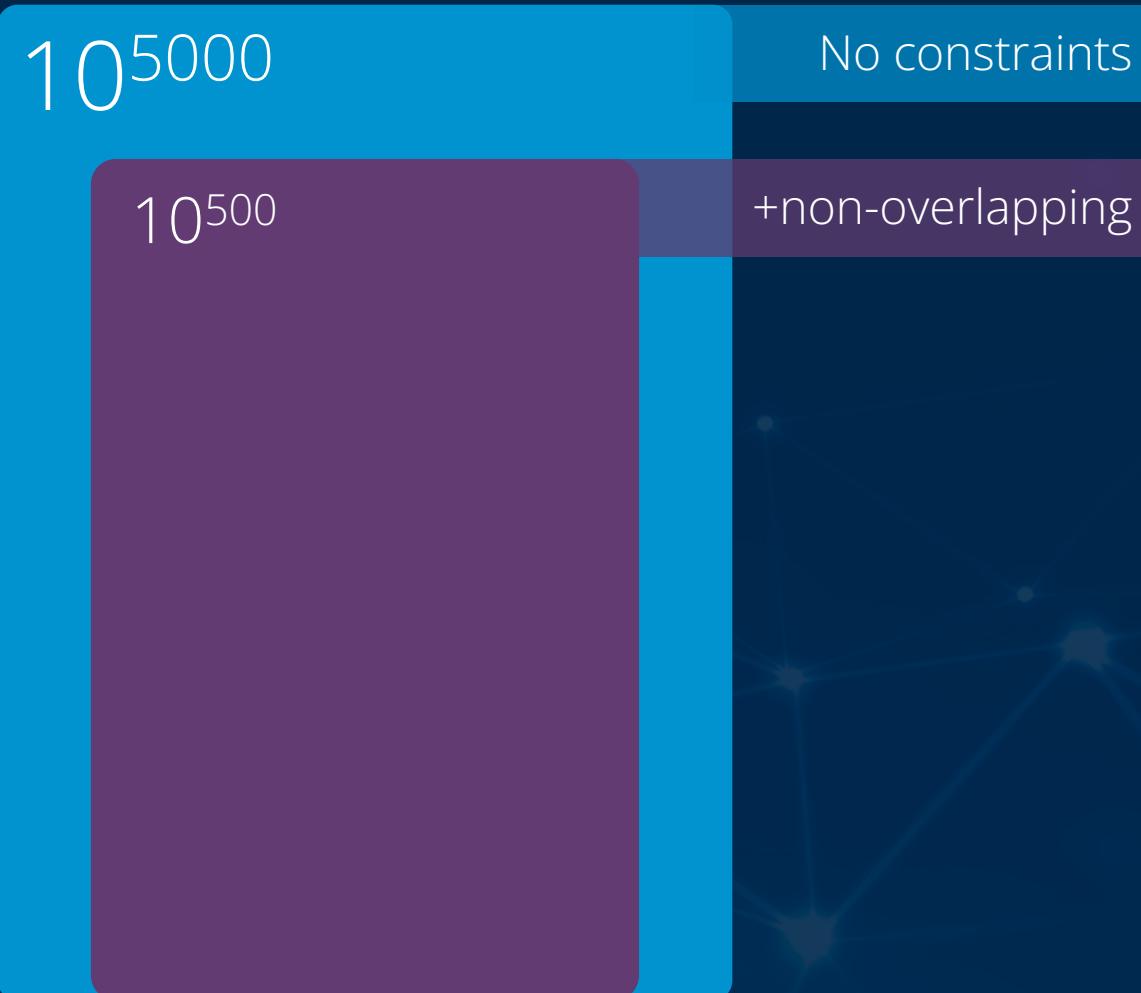
Solution space wrt. different constraints

Fixed-priority
(e.g., DM, RM)

FP+
offsets

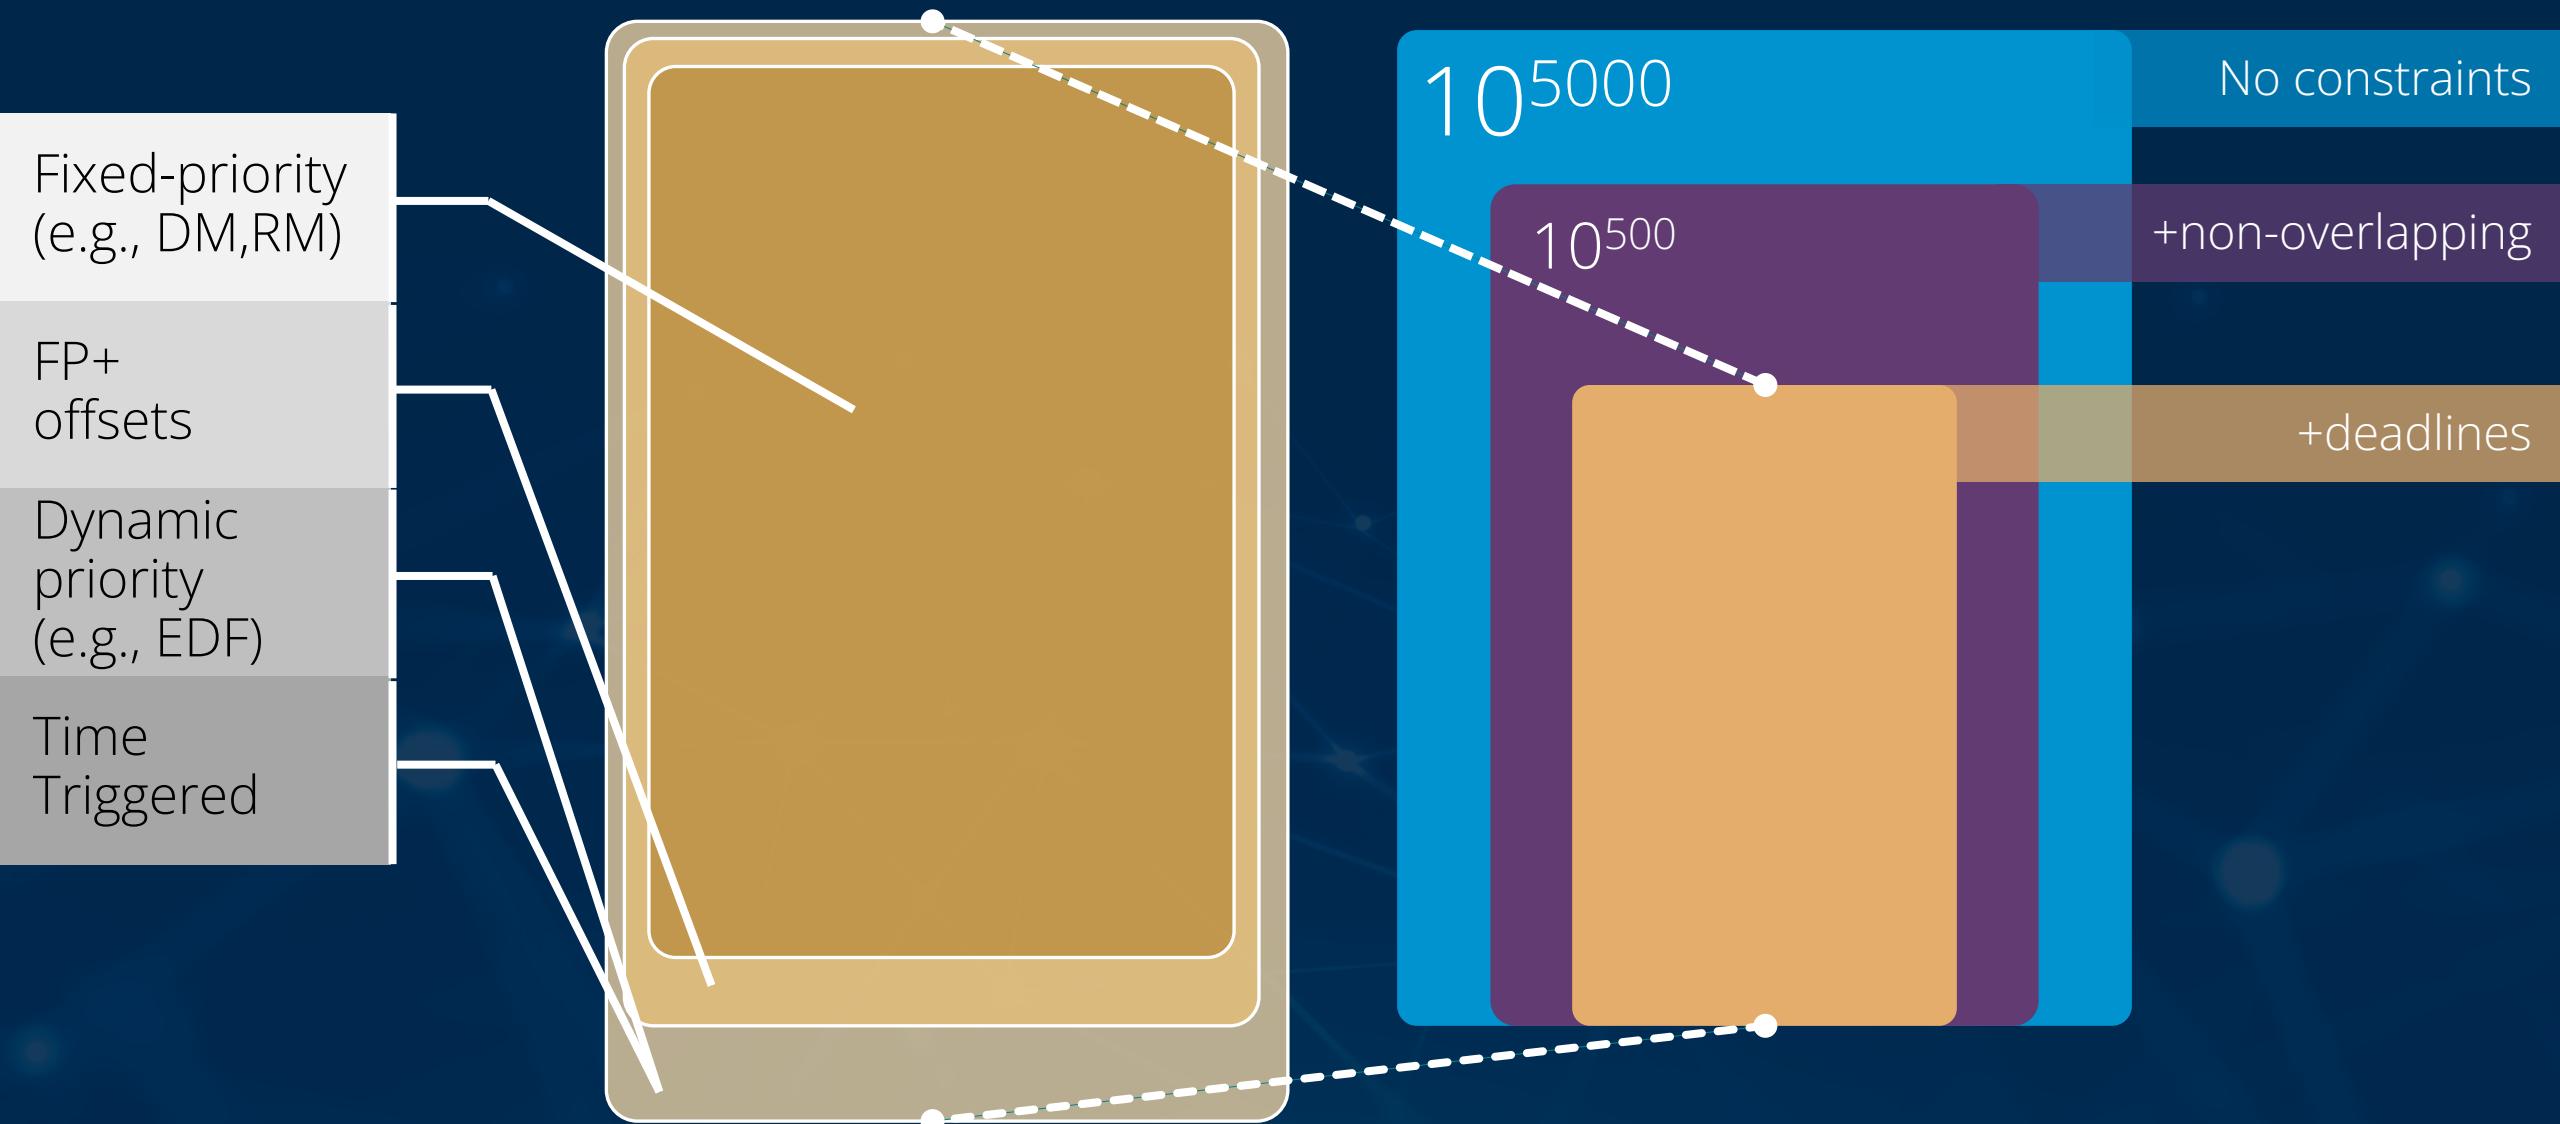
Dynamic
priority
(e.g., EDF)

Time
Triggered

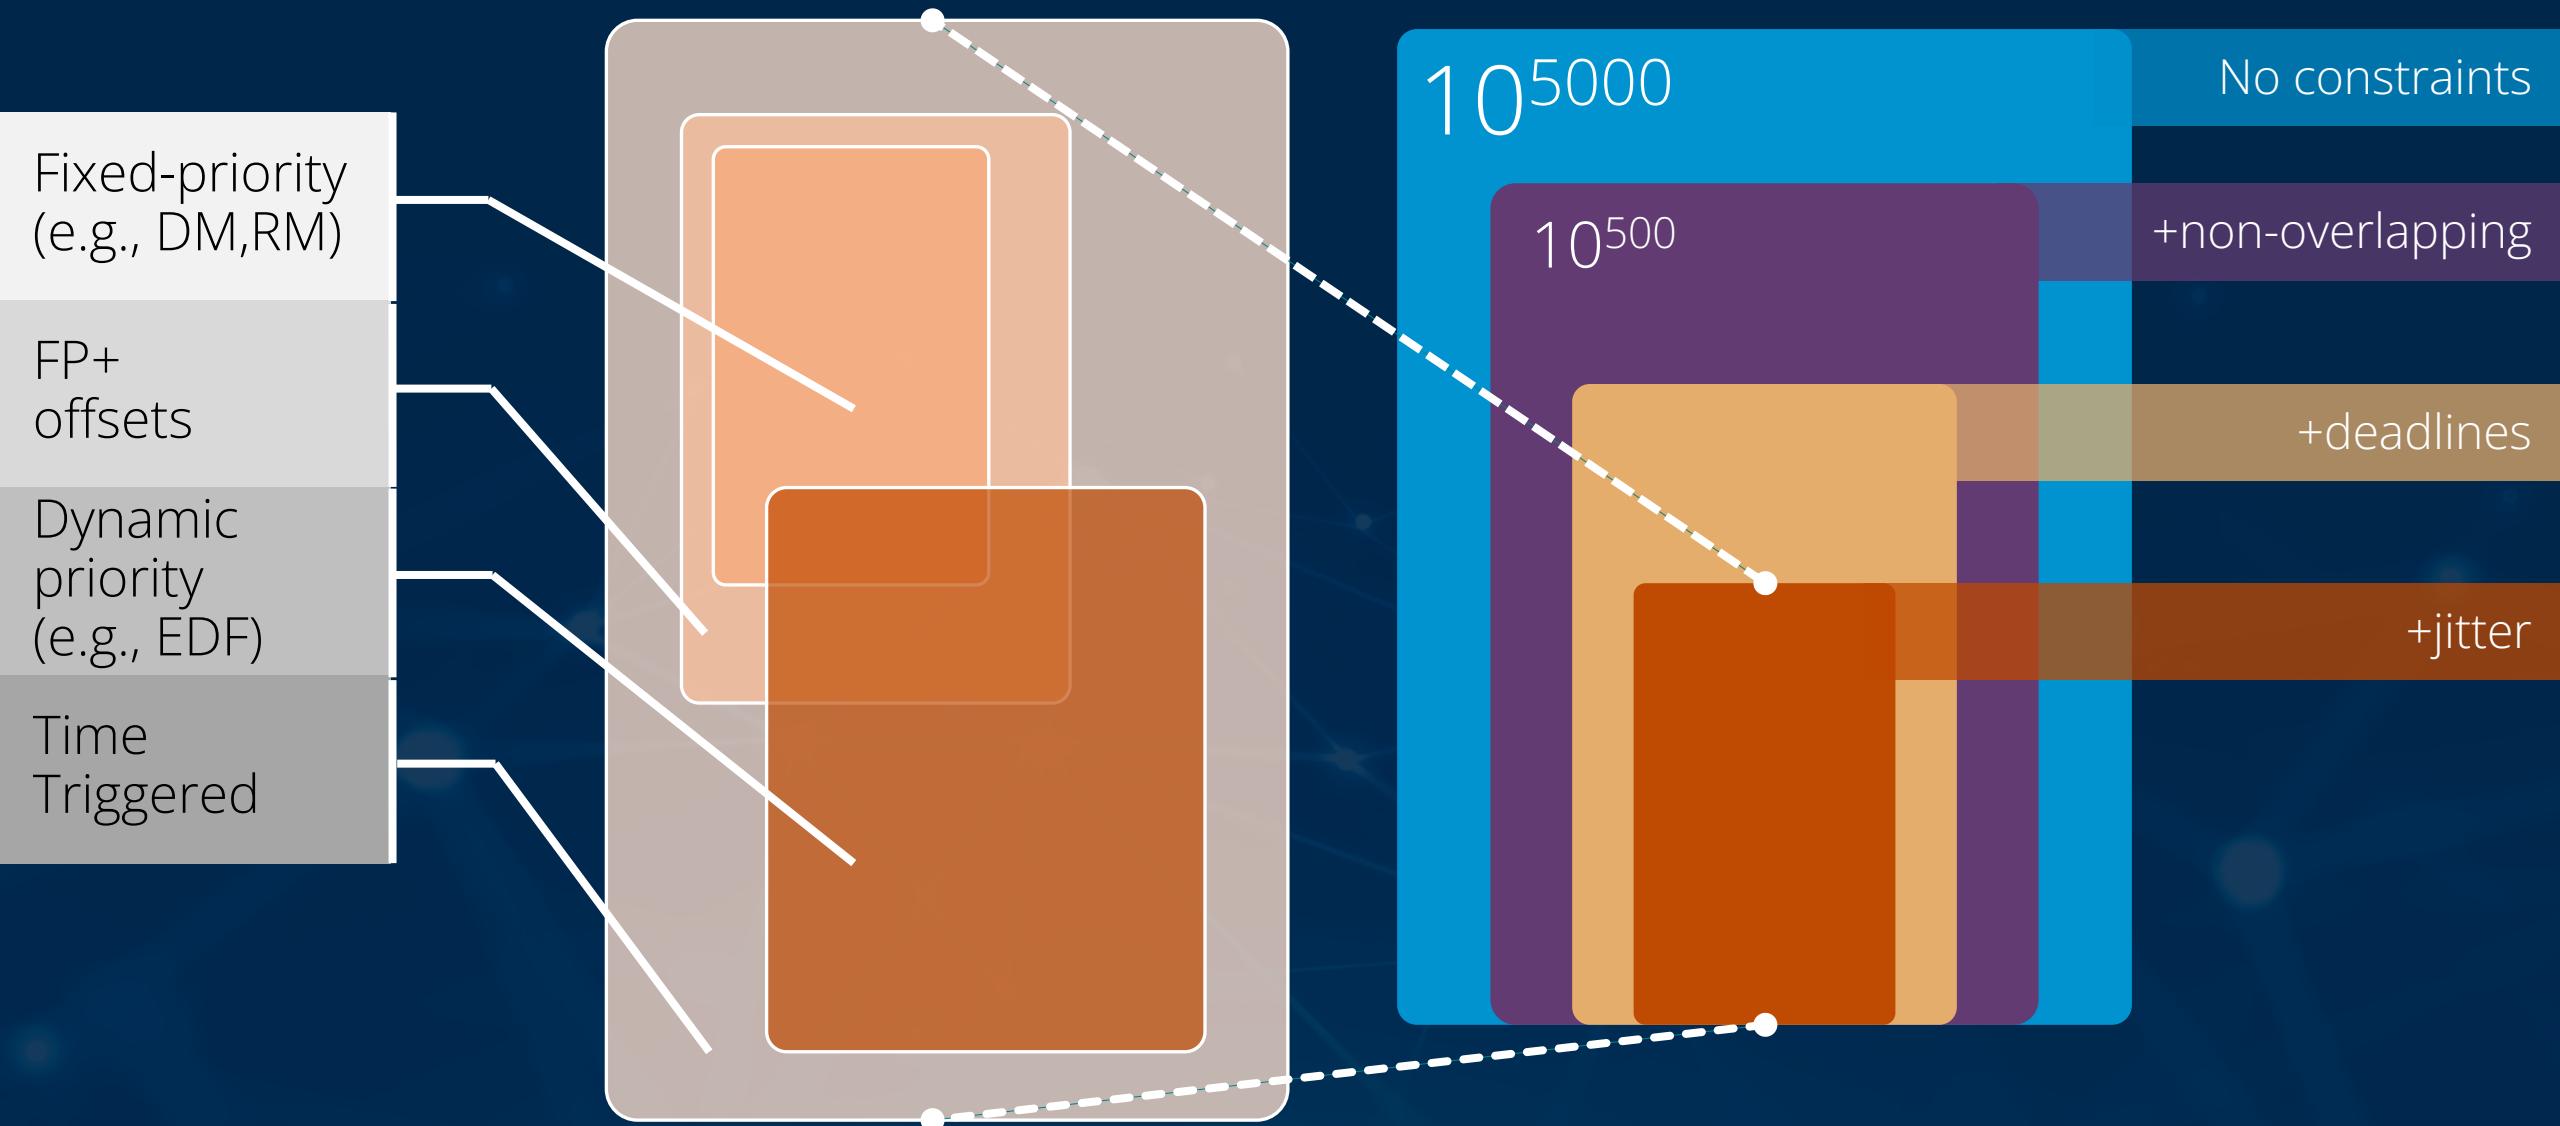


Only TT guarantees the exact position in time.
Other paradigms assign the position at runtime.

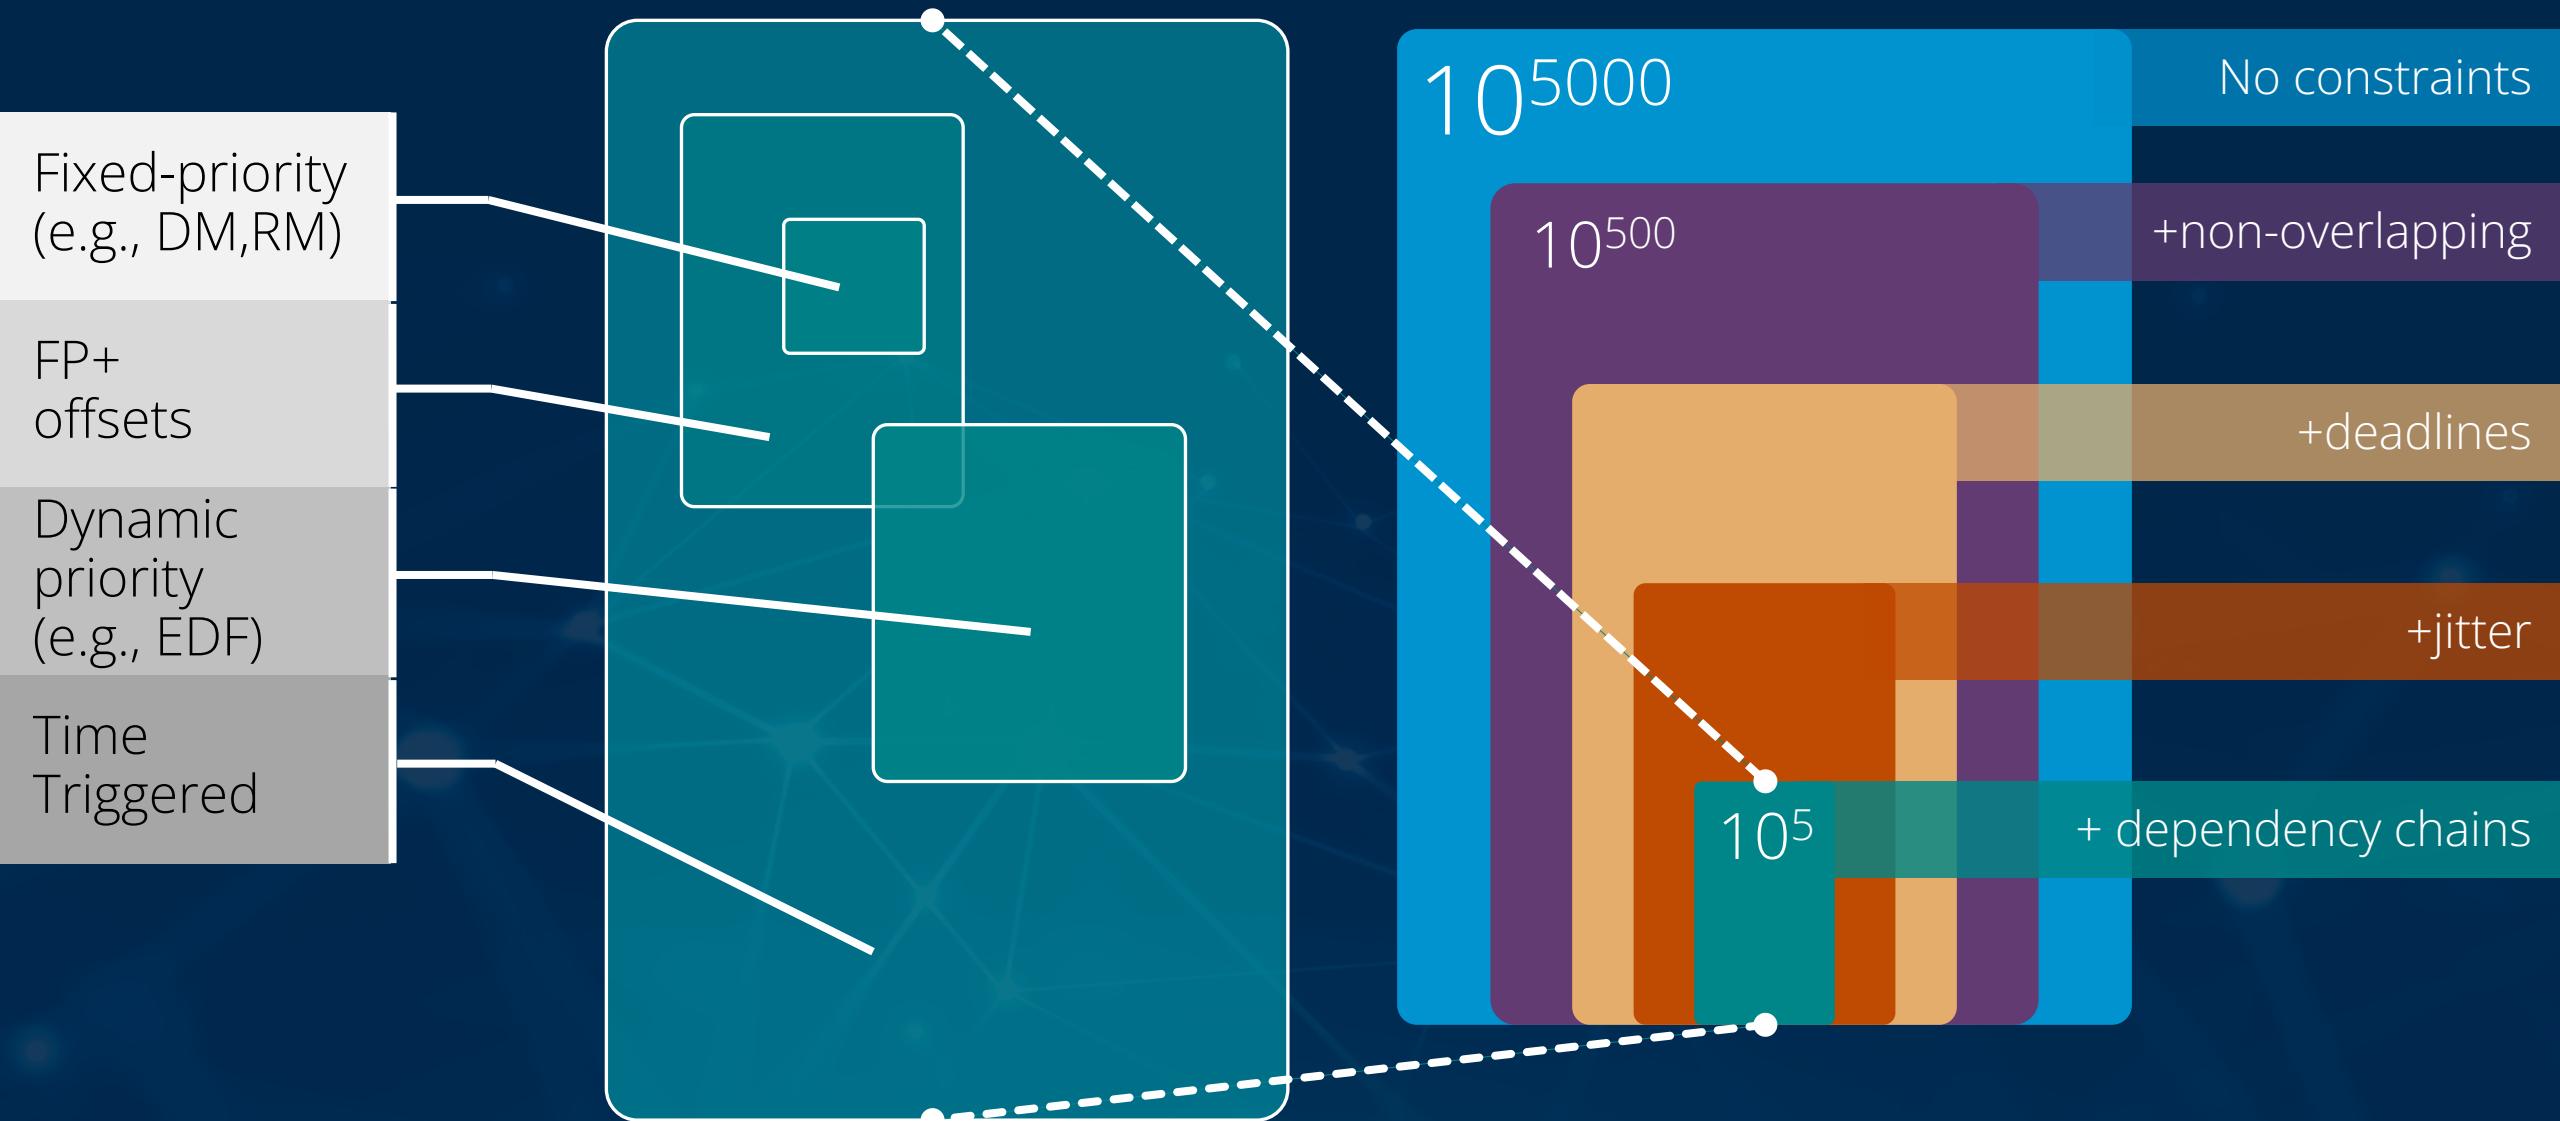
Solution space wrt. different constraints



Solution space wrt. different constraints



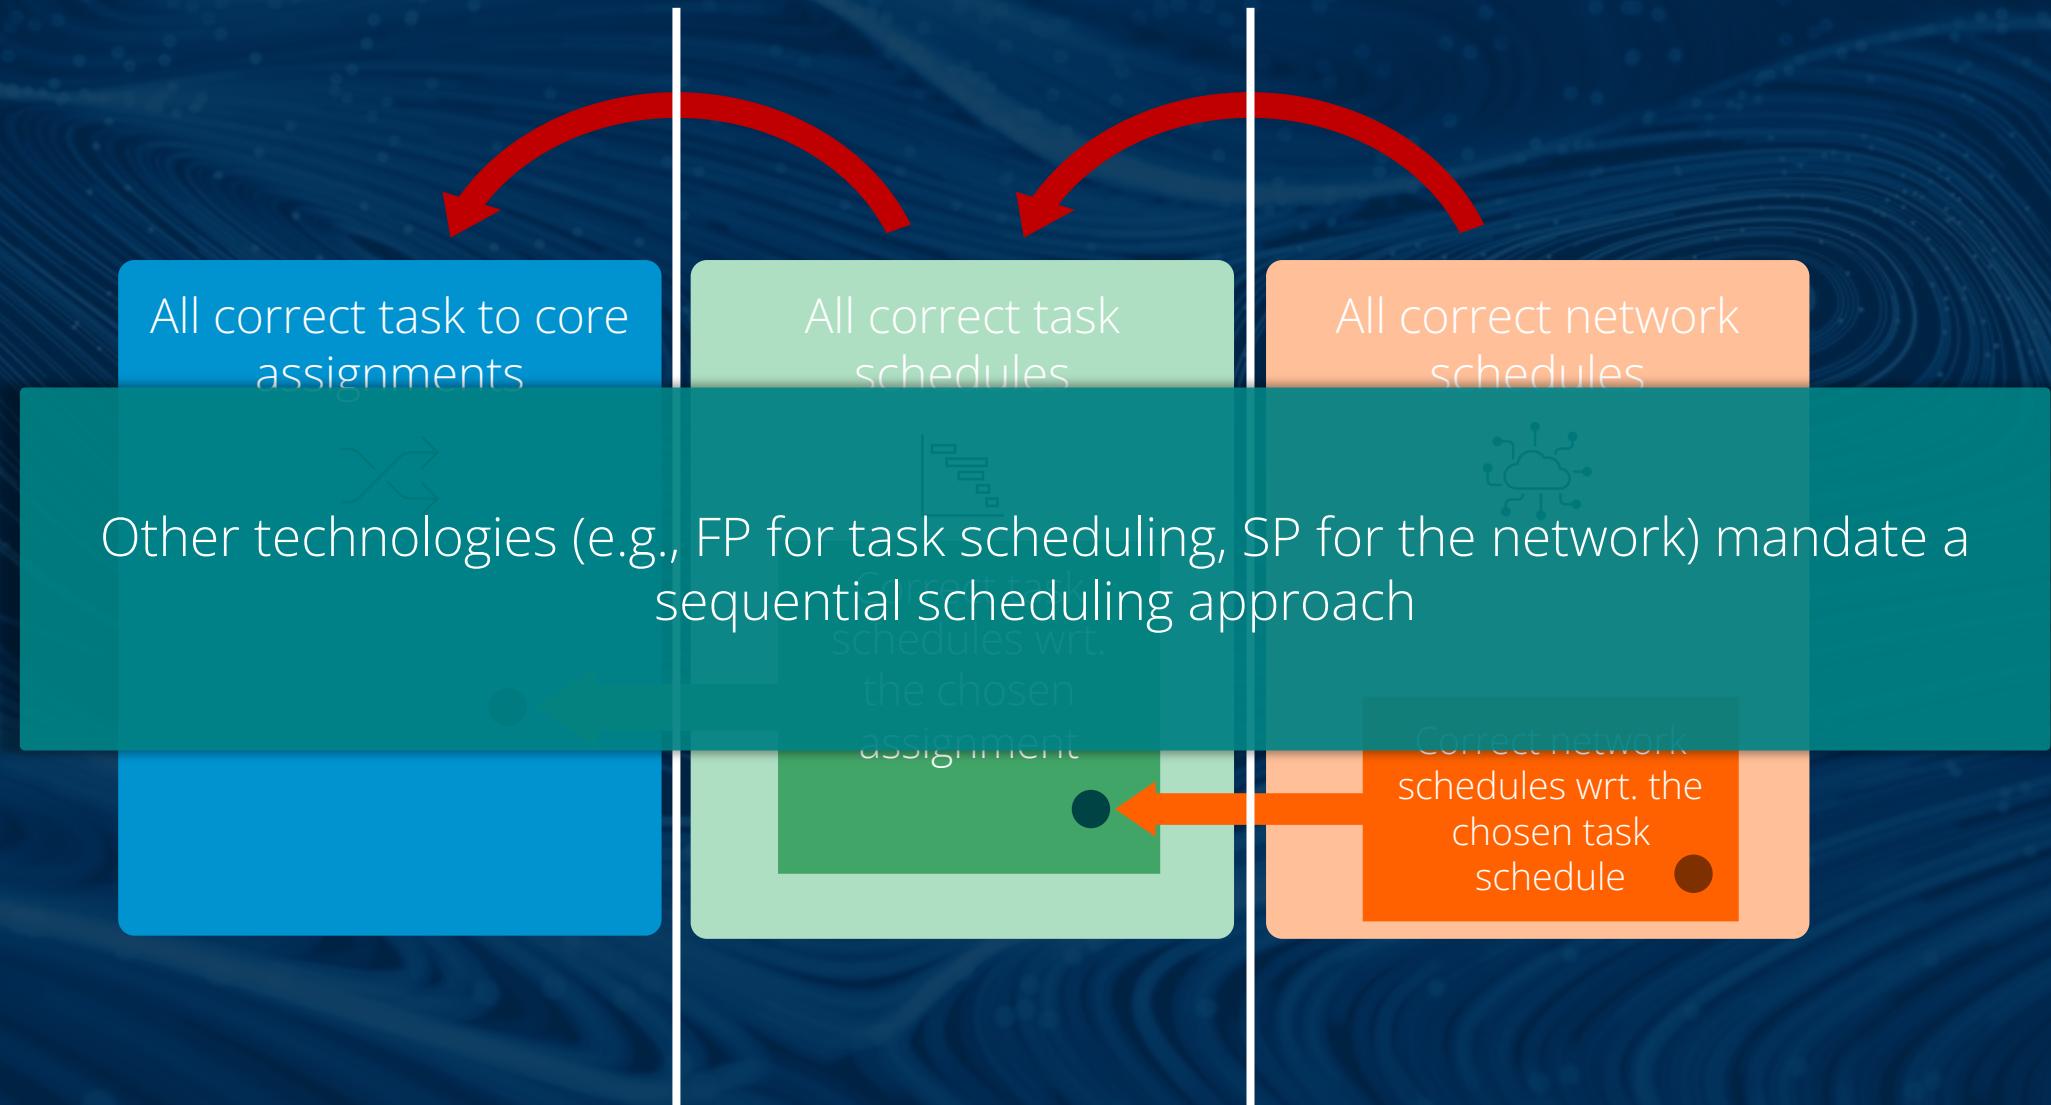
The time-triggered advantage



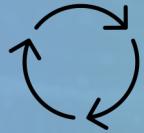
Problem dimensions



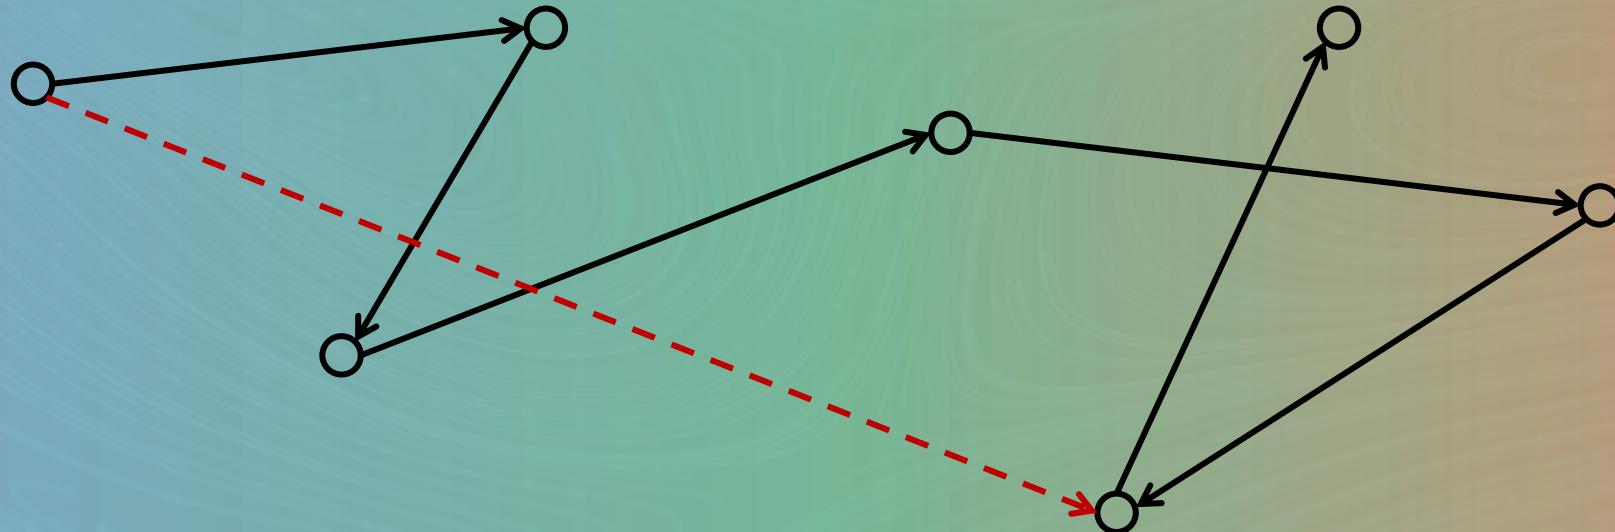
Sequential scheduling



Problem dimensions



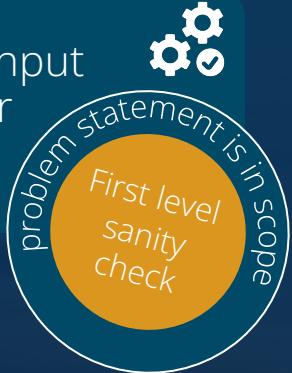
All possible correct schedules



Global Scheduling allows us to seamlessly go from one problem dimension to another, do fine-grained improvements, and even find relations between candidate solutions across the different dimensions

Global scheduler

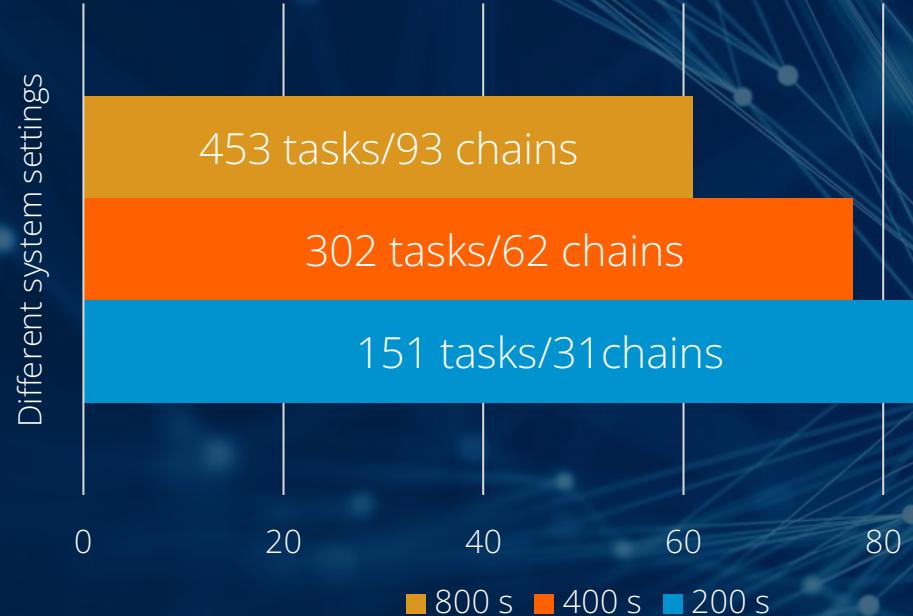
Formal Input
Compiler



Based on a rigorous formal model:

- non-interpretable correctness constraints expressed in mathematical language
- required for schedule correctness characterization and validation
- can be independently verified
- we have probably the most complete, rigorous, and comprehensive formal correctness model for task & network scheduling in the ADAS domain (~60 pages)

Smart heuristics



S.D. McLean, S.S. Craciunas, E.A. Juul Hansen, and P. Pop - Mapping and Scheduling Automotive Applications on ADAS Platforms using Metaheuristics. In Proc. ETFA, IEEE, 2020.





Beyond time-triggered

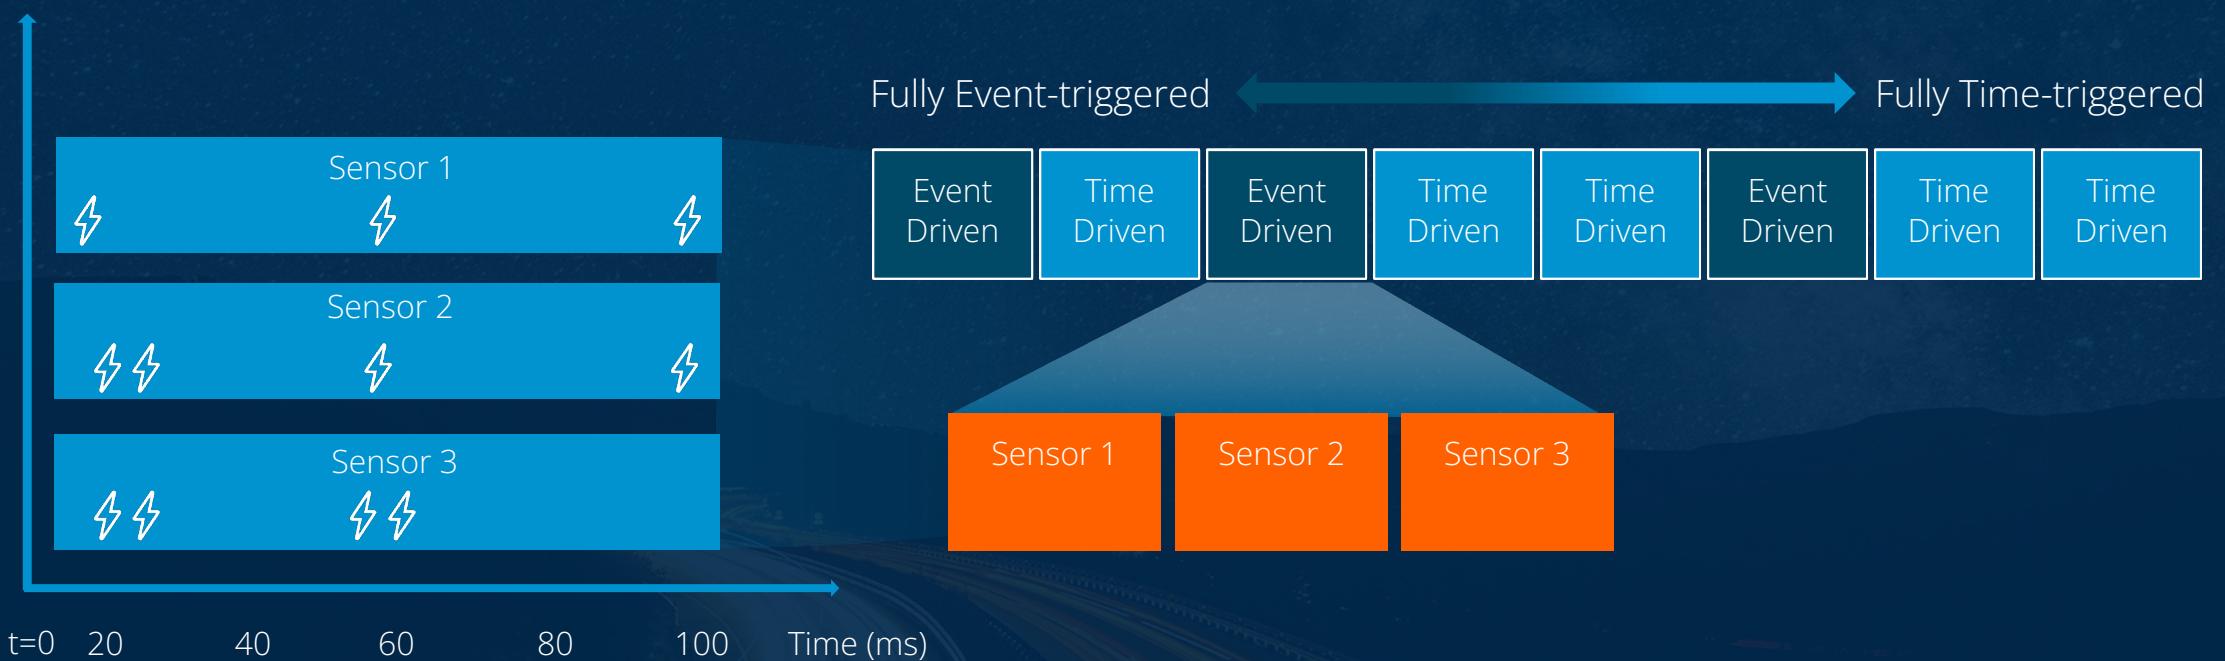
There is no free lunch!

- **Inflexible**: no (big) adaptations at runtime, task execution is fixed
 - event-triggered tasks (sporadic activation)
- **Overprovisioning**: Slots must be reserved for the worst-case
 - WCET analysis is hard and very pessimistic!
- **Precomputed**: task properties must be known at design time
 - this is a development process issue!
- **System limitations**: macrotick, hyperperiod
 - not really a problem in practical systems!
- **Event Driven tasks with guaranteed deadlines**
- Data-driven (DAG sequencer) execution in reserved slots
- Slot shifting to maximize CPU time and minimize timing budget estimation problems
- Scheduling Modes to switch between different schedules depending on driving modes



As much time-triggered as needed, as much flexibility as possible

Time-triggered and event/data-driven integration



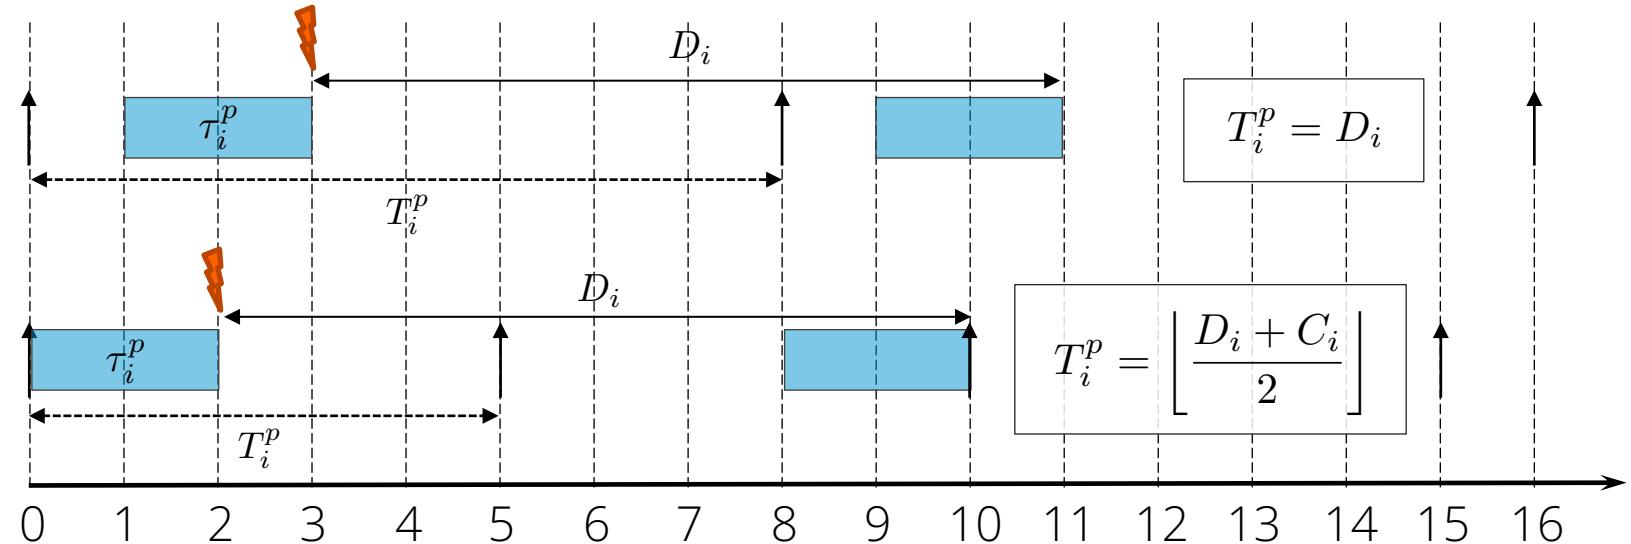
Second-level scheduling strategies based on requirements (e.g., FP)
Provide timing guarantees to critical event-driven tasks and optimize response times

How can we ensure that both TT and ET tasks respect their deadlines?

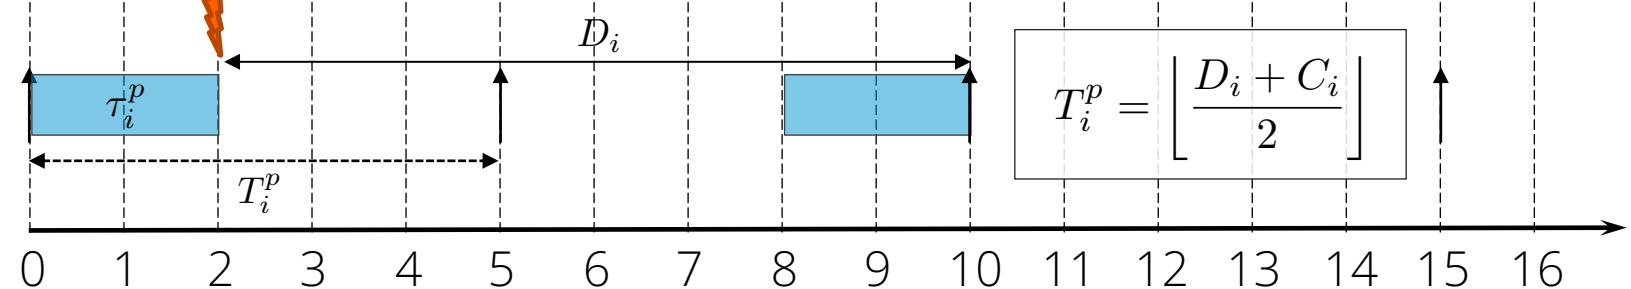
Oversampling – SPoll

Simple polling: Use one periodic polling task per event – oversampling
How long is the sampling period?

a) Strictly periodic



b) Non-strictly-periodic



Can be quite inefficient!

Think of an event with $C = 2$ ms, $D = 20$ ms, and $T = 100$ ms. (2% utilization)
We have to reserve a slot every 9 ms, consuming 22,2 % of CPU bandwidth.

Solution using response-time analysis

Holistic scheduling [Pop2003]

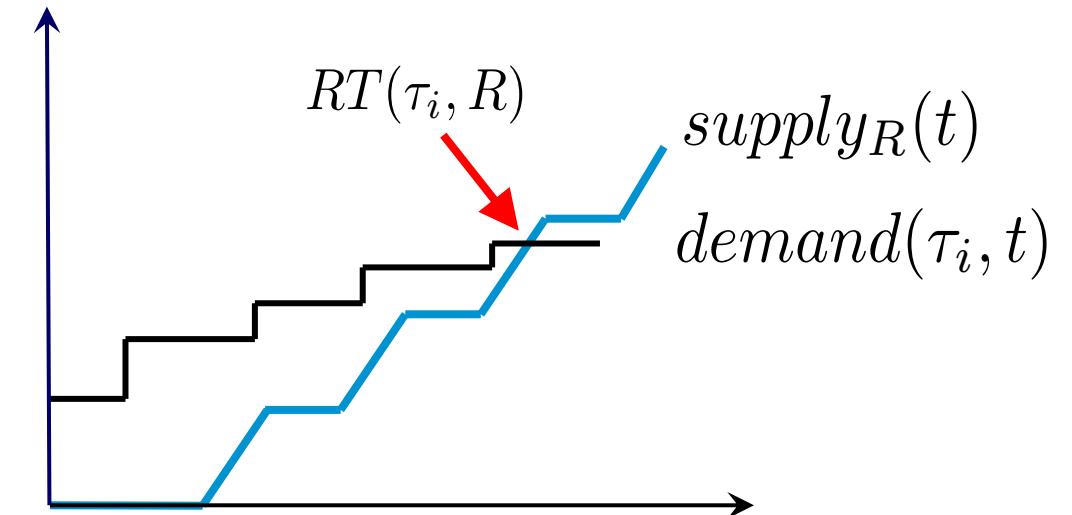
1. Generate TT schedule such that TT tasks are fulfilled.
2. Check schedulability over resulting idle slots using the response-time method over partitioned resources for every event-based task.
3. Recompute TT schedule if ET tasks not feasible

Slot-shifting [Isovic2009] is a similar method with a different schedulability test that does not assume FP scheduling.

Can be very time-consuming!
No guidance on how to place idle slots.



Can we do better?



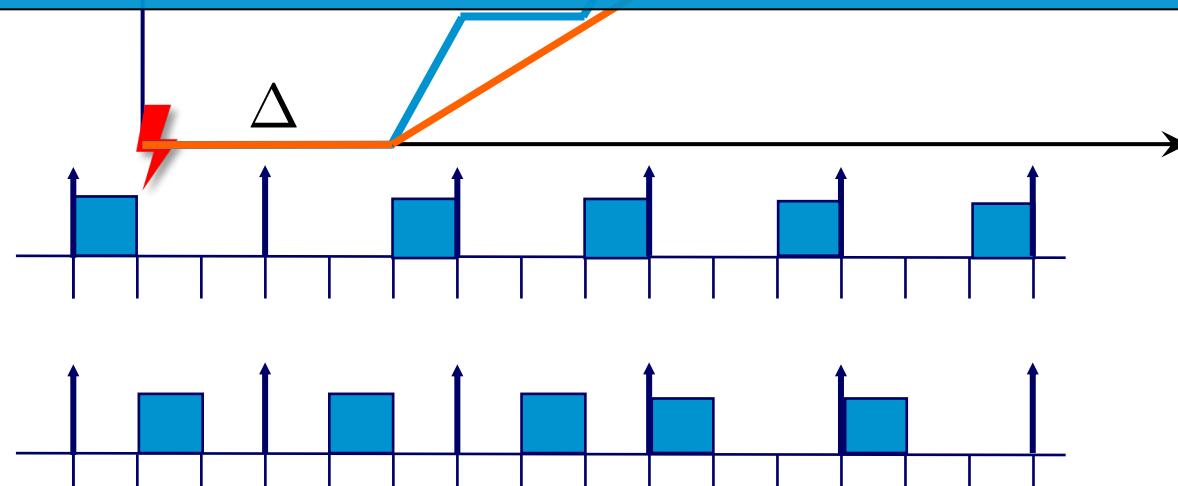
Solution using periodic resource abstraction

Explicit Deadline Periodic: R(C,D,T) [Shin2008]

- we do not care when the supply C is given, as long as it is given in every period T until the deadline D
- we can use the worst-case linear supply bound and schedulability test defined below [Almeida2004]
- New method: AdvPoll [Meroni2023]

- Schedulability analysis for ET is independent of the TT schedule
- The polling task (periodic resource) can be scheduled as a normal TT task

Tradeoff – schedulability for runtime



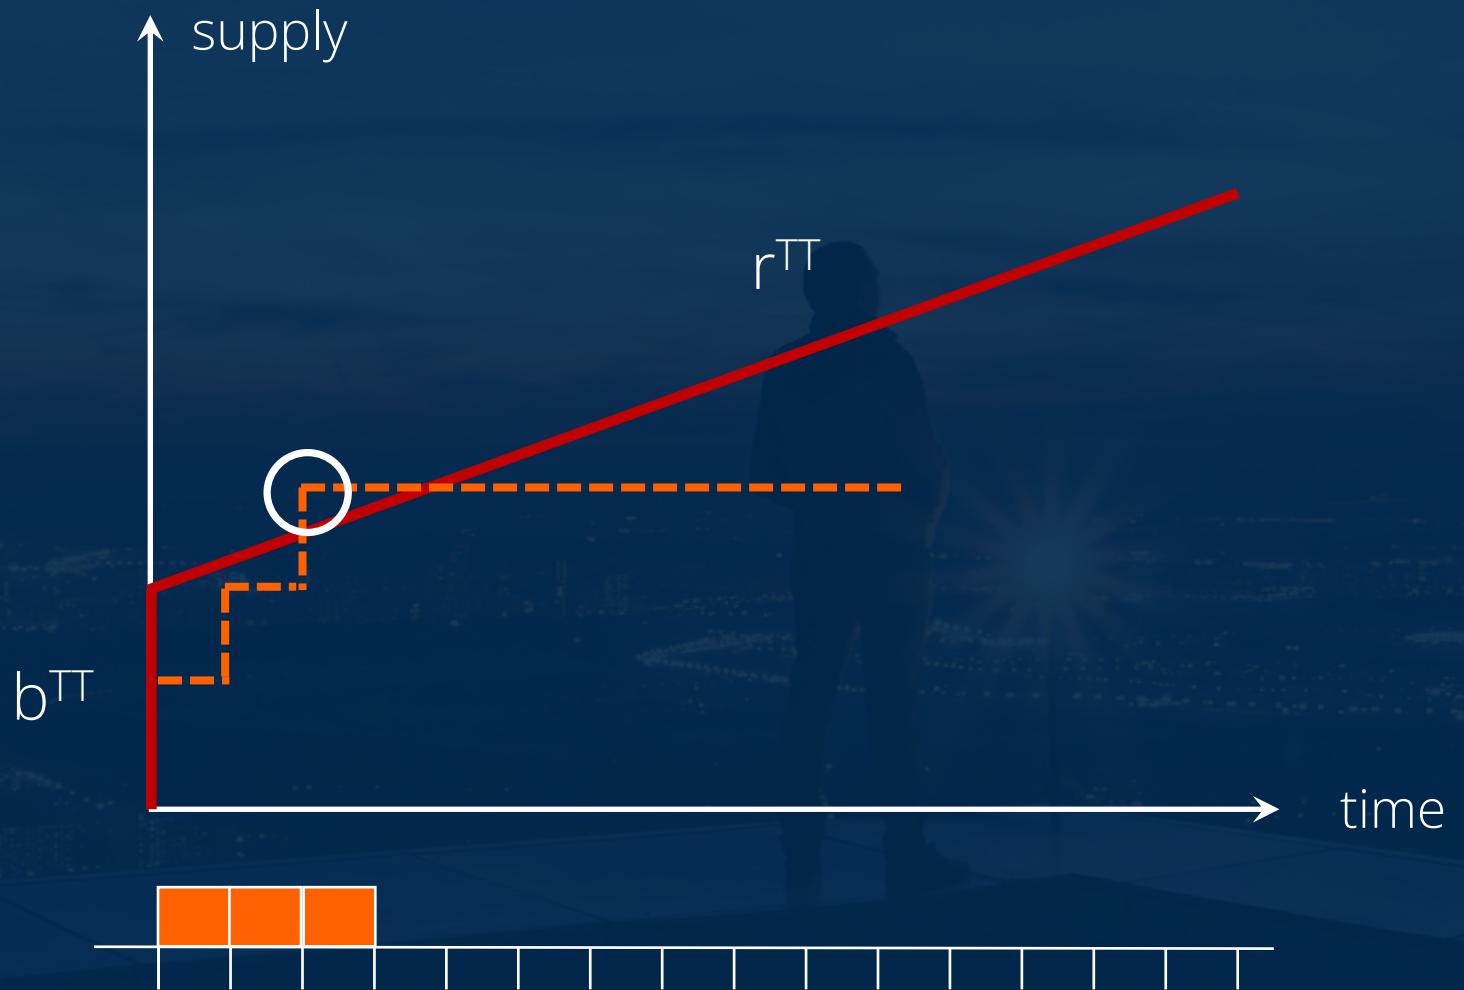
$$RT(\tau_i, R) = \text{earliest } t : t \geq \Delta + \text{demand}(\tau_i, t)/\alpha$$

$$RT(\tau_i, R) = \text{earliest } t : t \geq \Delta + \frac{1}{\alpha} \left(C_i + \sum_{\forall \tau_j \in HP(\tau_i)} \left\lceil \frac{t}{T_j} \right\rceil \cdot C_j \right)$$

Integration of ET + TT via real-time calculus

Main idea: derive a constraint on the TT schedule that will guarantee ET task schedulability & create a TT schedule that respects that constraint

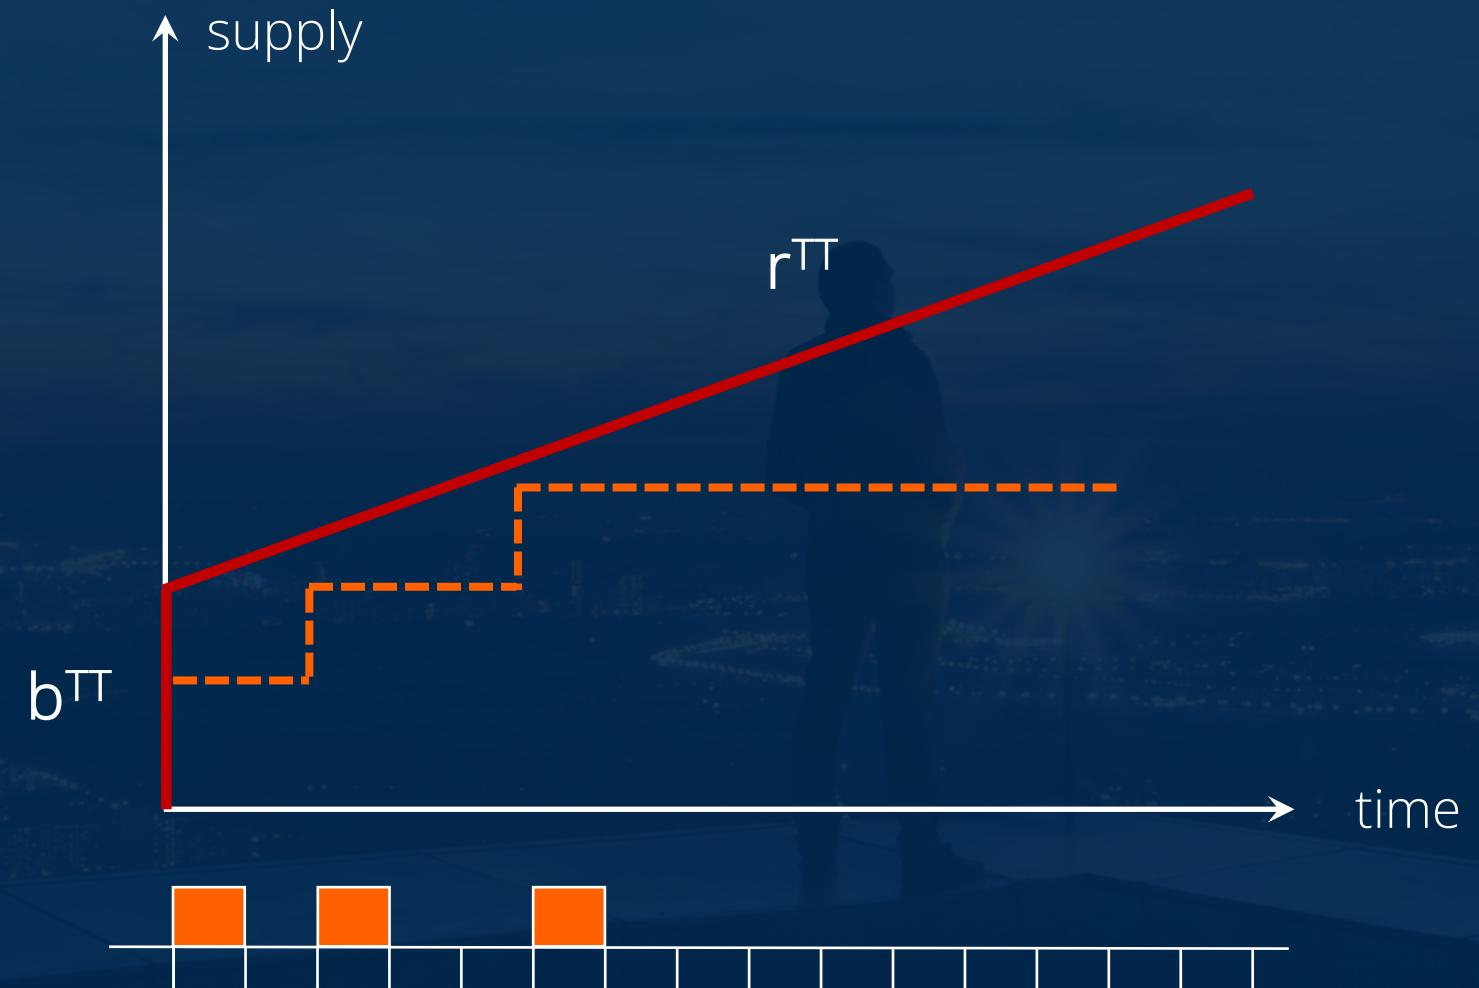
A. Finzi, S.S. Craciunas, M. Boyer –
Integrating Sporadic Events in Time-triggered Systems via Affine Envelope Approximations. In Proc. RTAS 2024



- o The cumulative computation time of TT must remain under the $(b^{\text{TT}}, r^{\text{TT}})$ constraint

Challenges:

- How to calculate the $(b^{\text{TT}}, r^{\text{TT}})$ constraint to enforce ET schedulability?
- How to compute an offline schedule remaining under the $(b^{\text{TT}}, r^{\text{TT}})$ constraint?



How to calculate the $(b^{\text{TT}}, r^{\text{TT}})$ constraint?

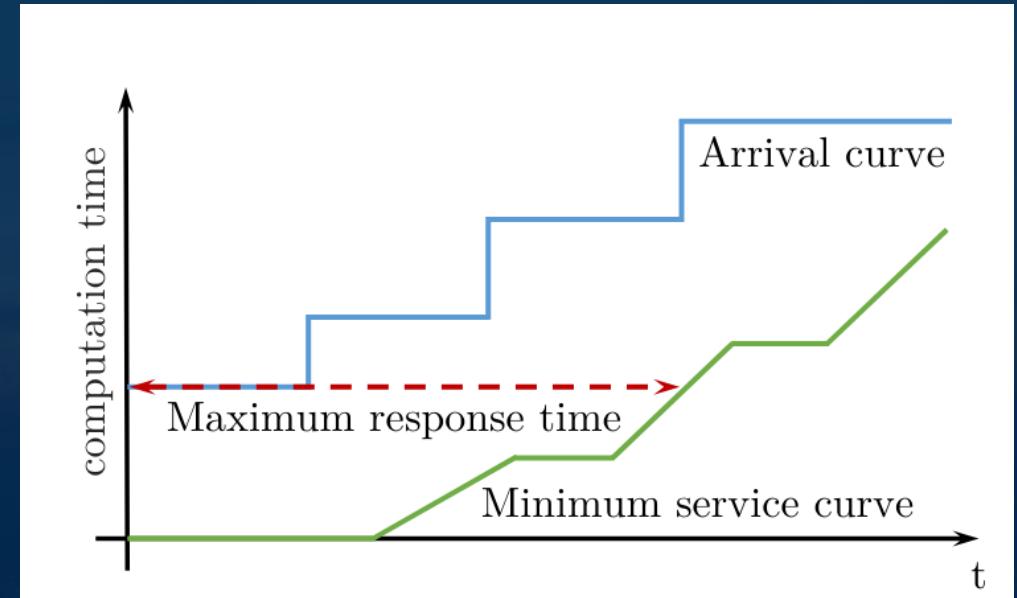
Real-Time Calculus:

Framework used to calculate the **maximum response time** of a task.

Arrival curves: describe how a set of tasks arrives in any time interval

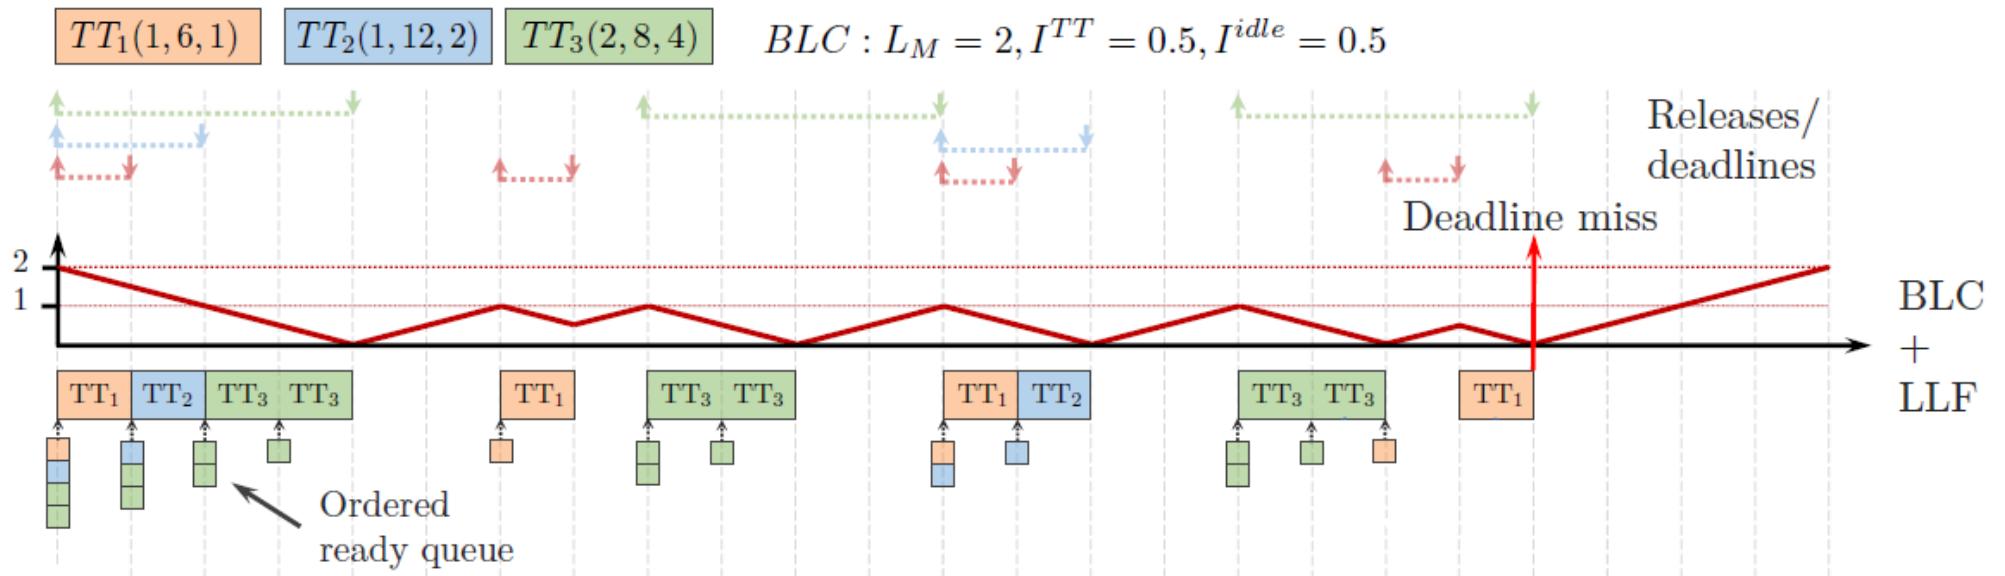
Service curves: describe the service offered by a CPU to a set of tasks in any time interval

- Maximum horizontal distance is the maximum response time (MRT)
- ET schedulability: $\text{MRT} \leq \text{deadline}$



- $(r^{\text{TT}}, b^{\text{TT}})$ are parameters of the ET service curves
- r^{TT} is the sum of the rates of the set of TT tasks
- The maximum b^{TT} such as the ET tasks are schedulable is computed in a binary search

How to compute a schedule remaining under the $(b^{\text{TT}}, r^{\text{TT}})$ constraint?



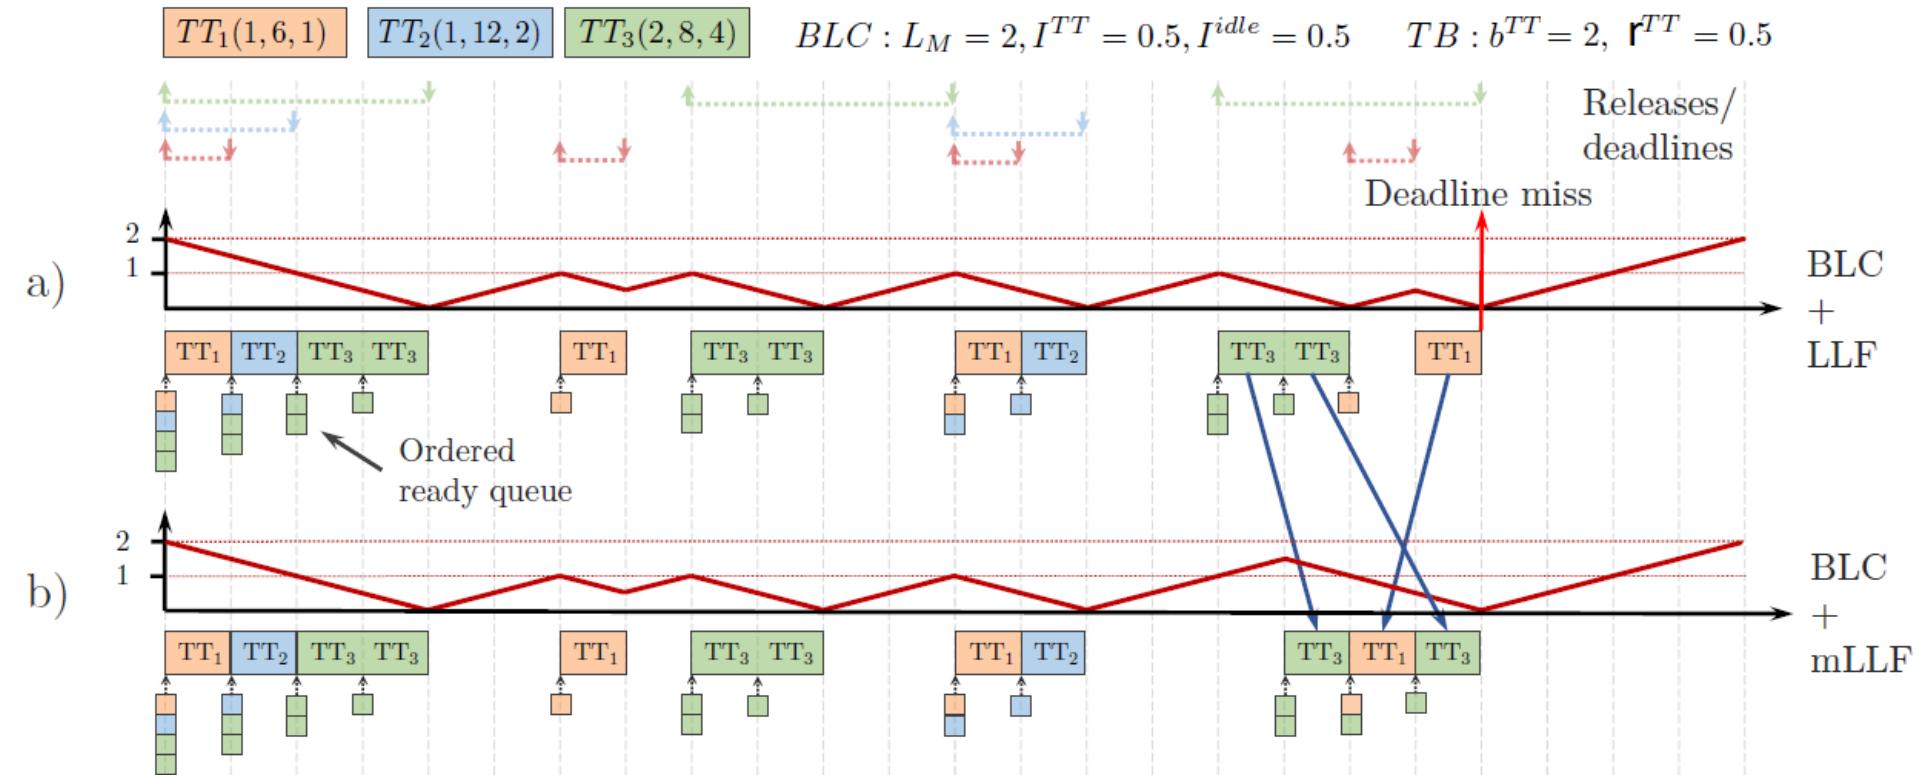
ET schedulability is enforced, but to the detriment of TT schedulability

- We need to balance ET and TT schedulability
 - We add “idle” tasks in the LLF, with Laxity:

$$L_{IDLE}(t) = \begin{cases} \left\lfloor \frac{bdg_{\sigma}(t)}{I^{TT}} \right\rfloor & \text{if } bdg_{\sigma}(t) < L_M - I^{\text{idle}} \\ \infty & \text{otherwise} \end{cases}$$

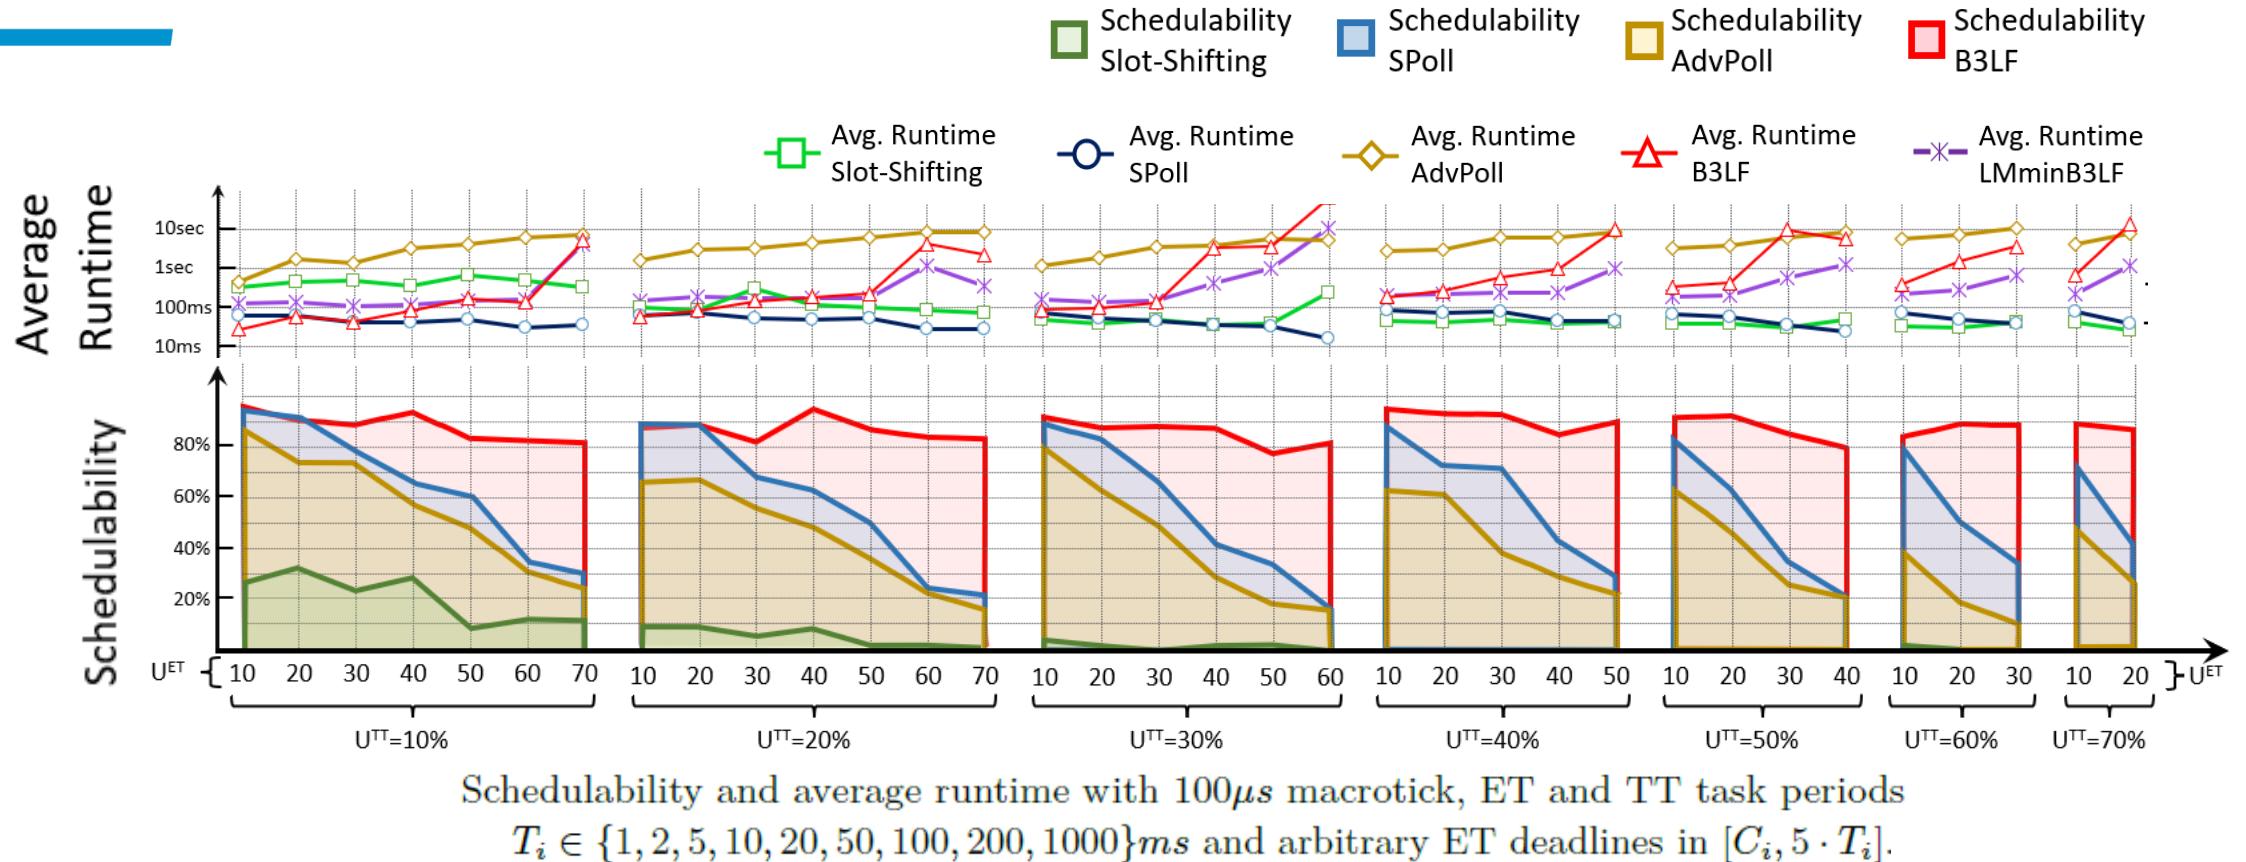
How to compute a schedule remaining under the (b^{TT}, r^{TT}) constraint?

With mLLF + BLC, we can find a valid TT schedule enforcing the ET schedulability



Experiment results:

With single core CPU (or a core in a fully partitioned multicore system)
 U^{TT} and U^{ET} are the respective utilizations of the TT and ET task sets



Schedulability of ET and TT tasks with B3LF is better than the other solutions

Thank you!

[Xu2000] Xu, Parnas - Priority Scheduling Versus Pre-Run-Time Scheduling. Real-Time Syst, 2000

[Pop2003] T. Pop, P. Eles, Z. Peng - Schedulability analysis for distributed heterogeneous time/event triggered real-time systems, ECRTS 2003.

[Isovic2009] D. Isovic, G. Fohler - Handling mixed sets of tasks in combined offline and online scheduled real-time systems, Real-Time Syst, 2009.

[Shin2008] I. Shin et al.- Hierarchical Scheduling Framework for Virtual Clustering of Multiprocessors. ECRTS, 2008

[Almeida2004] L. Almeida, P. Pedreiras - Scheduling within temporal partitions: Response-time analysis and server design. EMSOFT, 2004.

[Meroni2023] C. Meroni, S.S. Craciunas, A. Finzi, and P. Pop - Mapping and Integration of Event- and Time-triggered Real-time Tasks on Partitioned Multi-core Systems. In Proc. ETFA, 2023.

[Finzi2024] A. Finzi, S.S. Craciunas, M. Boyer – Integrating Sporadic Events in Time-triggered Systems via Affine Envelope Approximations. In Proc. RTAS, 2024

We are hiring!



<https://www.tttech.com/jobs-career/>